

Unsupervised Learning and Auto-Encoders

Neural Computation Course 2022

Konstantinos Kamnitsas

Intro to Unsupervised Learning


Neural Computation Course 2022

Konstantinos Kamnitsas


Previously: Supervised Learning

Goal: learn a function to map $f_{\theta} : \mathcal{X} \rightarrow \mathcal{Y}$

Given for training:
Input & label $(x_1, y_1), \dots, (x_N, y_N) \sim p_{data}(x, y)$




Regression




Classification




Segmentation



Creating labelled data is challenging



Creating labelled data is challenging





Labelled data are limited

Creating labelled data requires:

- Time consuming manual work
- Humans are expensive
- Labelling may require expertise (e.g. healthcare)

Consequence: Labelled databases are small



A DAY IN DATA

The exponential growth of data is undisputed, but the numbers behind this explosion – fuelled by internet and the use of connected devices – are hard to comprehend, particularly when looked at in the context of one day.

500m

tweets are sent every day

Twitter

294bn

billion emails are sent

Radicati Group

320bn

emails to be sent each day by 2021

306bn

emails to be sent each day by 2020

3.9bn

people use emails

4PB

of data created by Facebook, including

350m photos

100m hours of video watch time

Facebook Research

4TB

of data produced by a connected car

Intel

Unlabeled data is in abundance!

DATA UNITS		
TB	terabyte	1,000 ³ bytes
PB	petabyte	1,000 ³ TB
EB	exabyte	1,000 ³ PB
ZB	zettabyte	1,000 ³ EB
YB	yottabyte	1,000 ³ ZB

*A lowercase "b" is used as an abbreviation for bits, while an uppercase "B" represents bytes.

65bn

messages sent over WhatsApp and two billion minutes of voice and video calls made

Facebook

95m

photos and videos are shared on Instagram

Instagram Business

28PB

to be generated from wearable devices by 2020

Statista

ACCUMULATED DIGITAL UNIVERSE OF DATA

4.4ZB

2013

44ZB

2020

PwC

Searches made a day

5bn

Searches made a day from Google

3.5bn

Smart Insights

From: <https://www.raconteur.net/infographics/a-day-in-data/>

Unsupervised Learning

Available data: $x_1, \dots, x_N \sim p_{data}(x)$

Goal: Learn “**useful features**” of the data / Learn “**the structure**” of the data.



Unsupervised Learning

Available data: $x_1, \dots, x_N \sim p_{data}(x)$

Goal: Learn “**useful features**” of the data / Learn “**the structure**” of the data.


Useful for:

- Dimensionality Reduction (compression)
- Clustering
- Probability Density Estimation
- Generation / Synthesis
- Learn from loads unlabeled data, when labeled data are limited
- ...


Generation / Synthesis



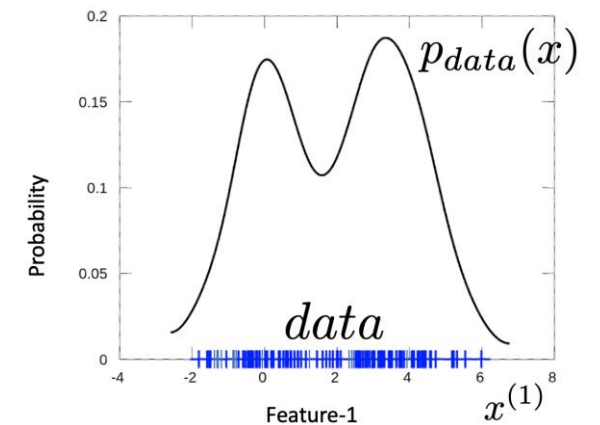
Dimensionality Reduction



Clustering



Probability Density Estimation

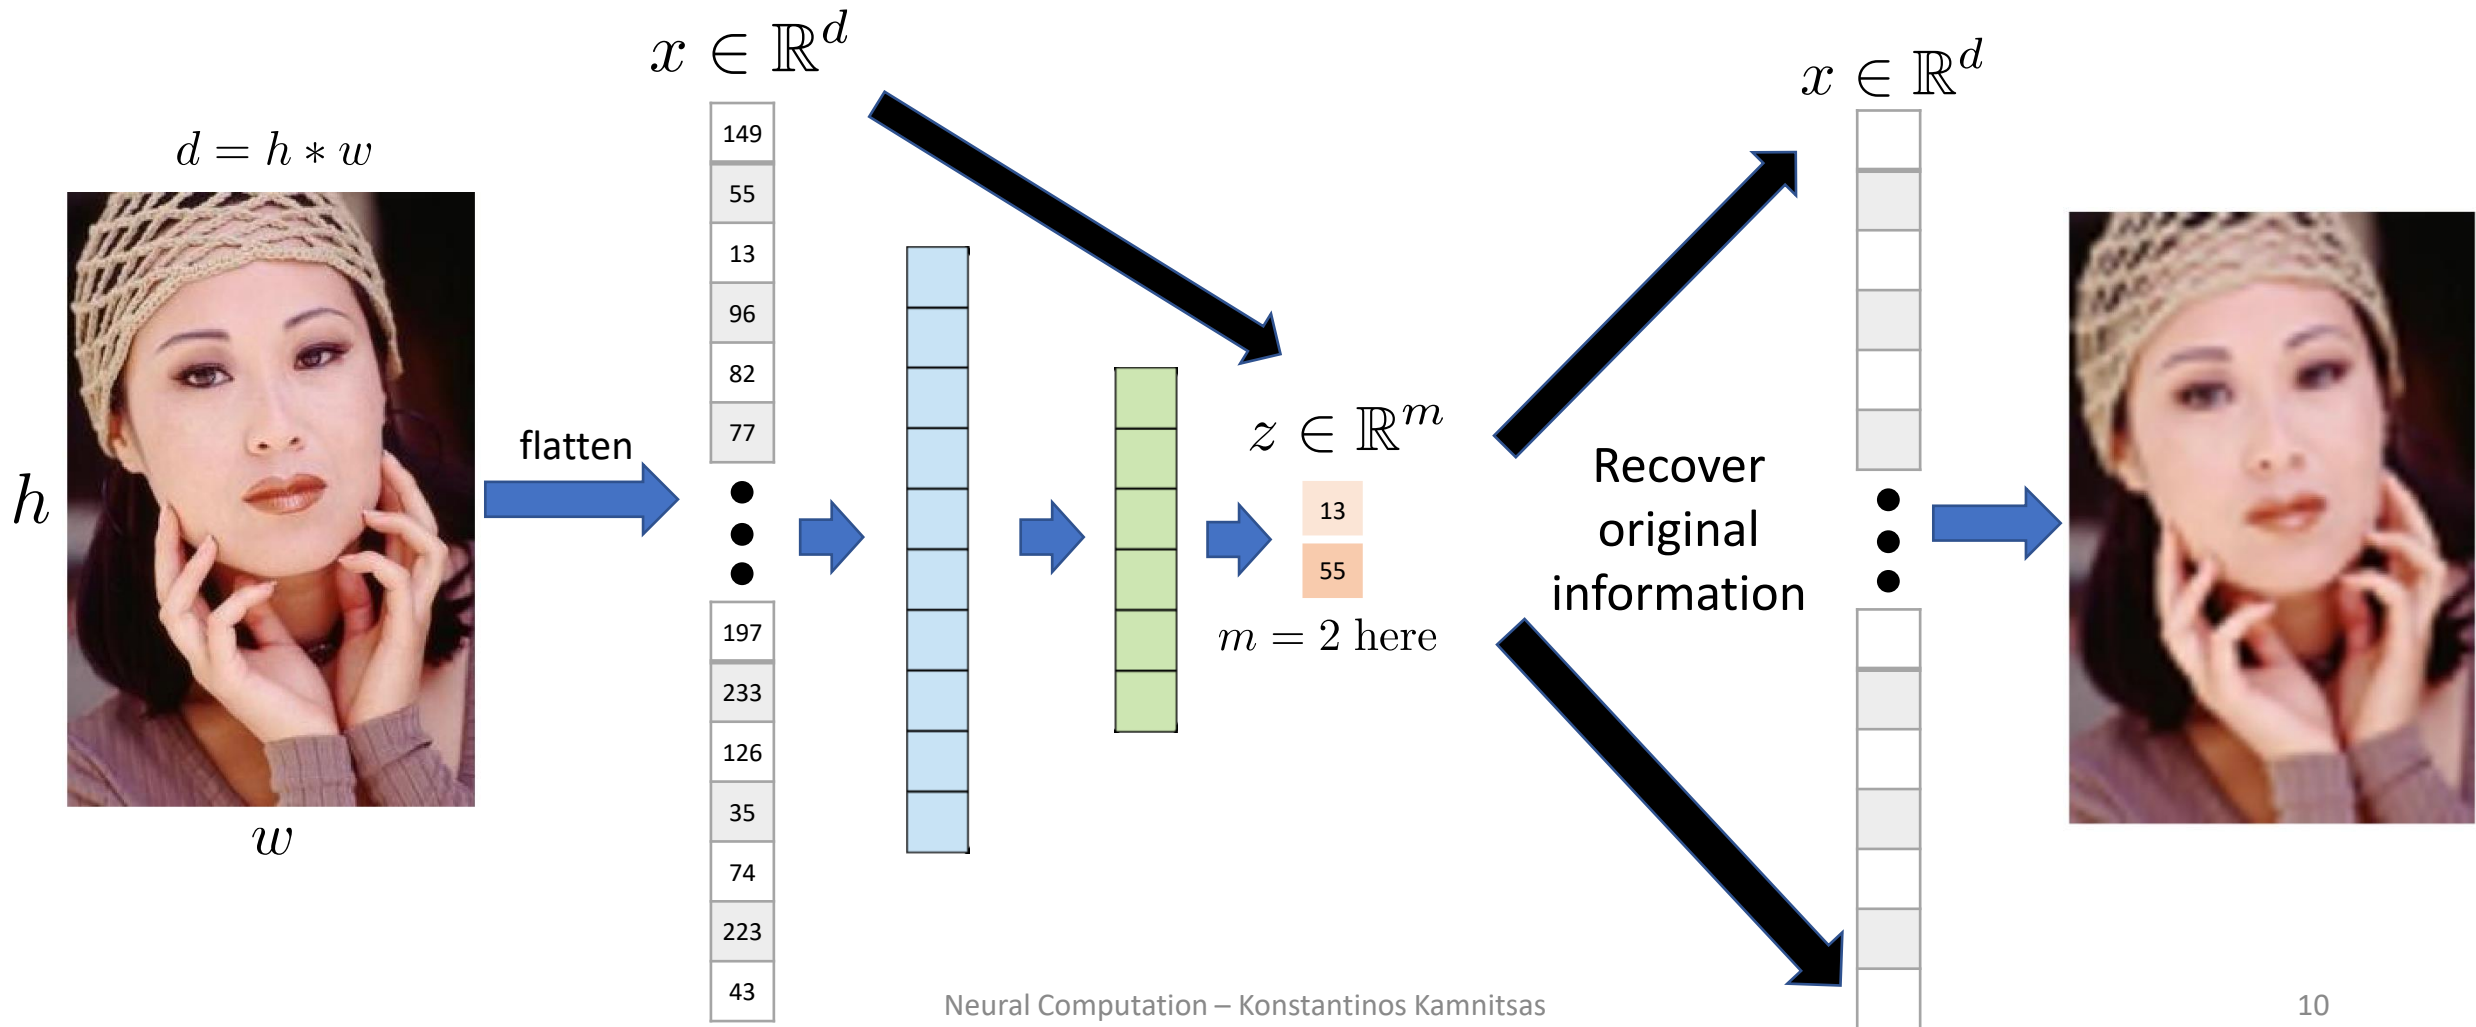


Dimensionality Reduction / Compression

Learn $f : \mathbb{R}^d \rightarrow \mathbb{R}^m$, where $d > m$

Requirement:

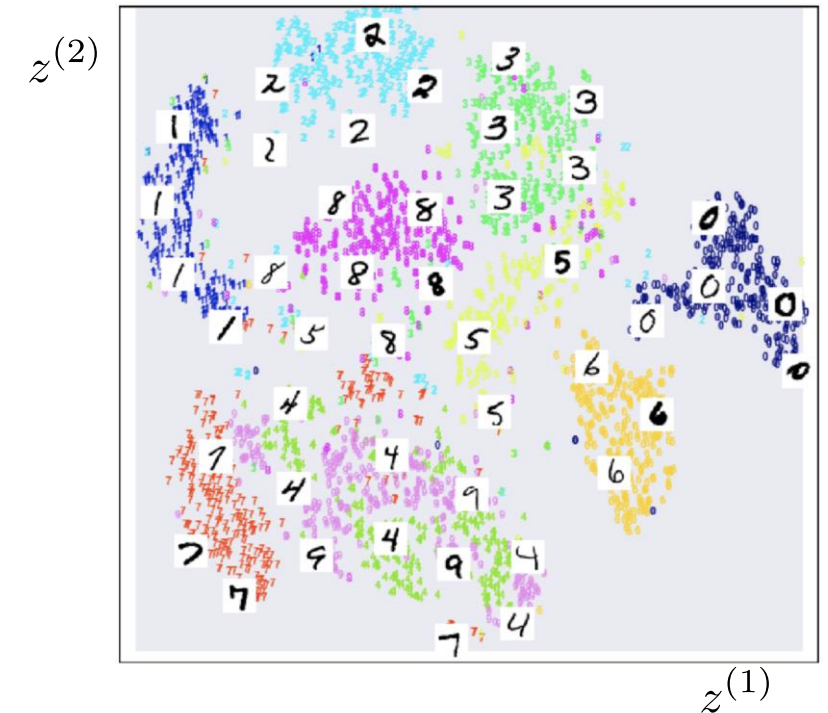

Preserve important information



Clustering

What is clustering: Discover or form groups of samples based on their similarity. Similar samples should be grouped together, dissimilar samples should be separated.

Example: Assume 10 “classes” of digits. Learn a function (model) that embeds (maps) samples of the same class to similar values (features), whereas samples of different classes are embedded to different values (features)



Clustering

What is clustering: Discover or form groups of samples based on their similarity. Similar samples should be grouped together, dissimilar samples should be separated.

Example: Assume 10 “classes” of digits. Learn a function (model) that embeds (maps) samples of the same class to similar values (features), whereas samples of different classes are embedded to different values (features)


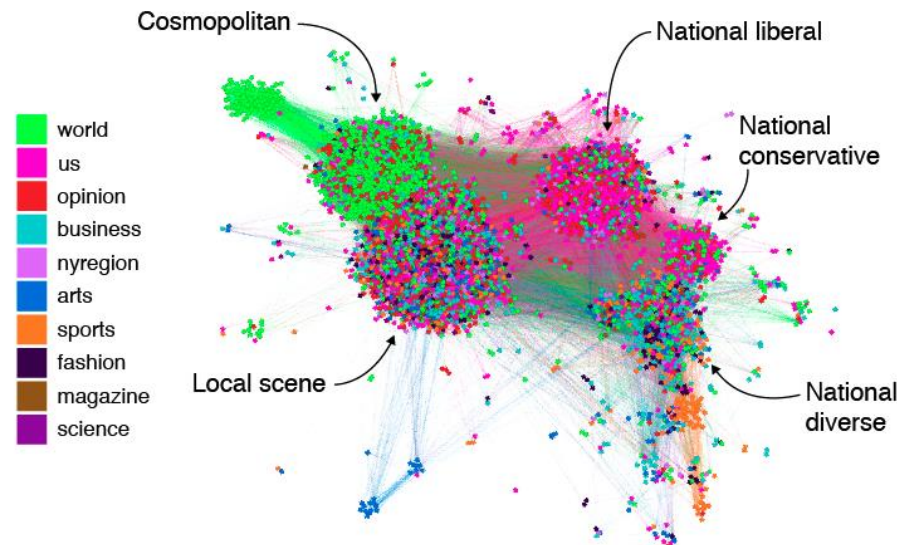
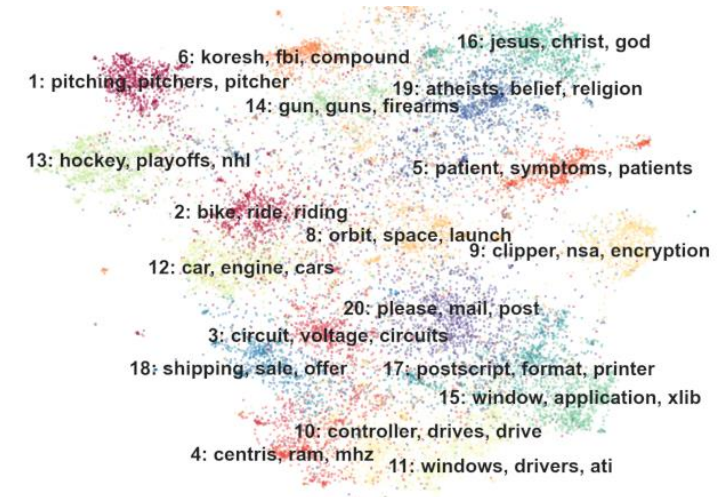


Image (digit) analysis



From: www.zdnet.com/article/new-study-maps-social-identity-clusters-on-twitter/


Social media users



Text analysis


Learn from **unlabeled data** to improve Supervised Learning **when labelled data are limited**

Unsupervised learning of clusters




Followed by

Supervised learning with few labeled data




Probability Density Estimation

1-dimensional data:

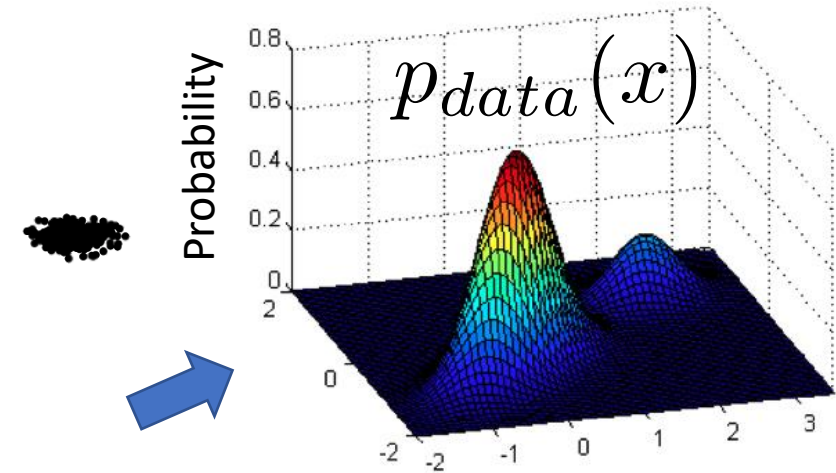
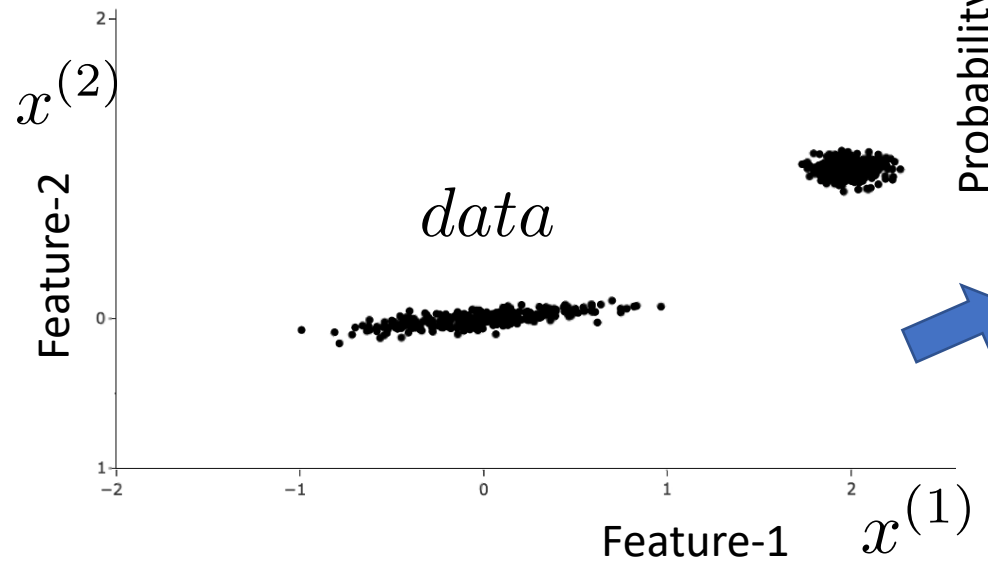


Probability Density Estimation

1-dimensional data:



2-dimensional data:



Generation / Synthesis

<https://thispersondoesnotexist.com/>



Generation / Synthesis

State of the Art 2014



State of the Art 2018



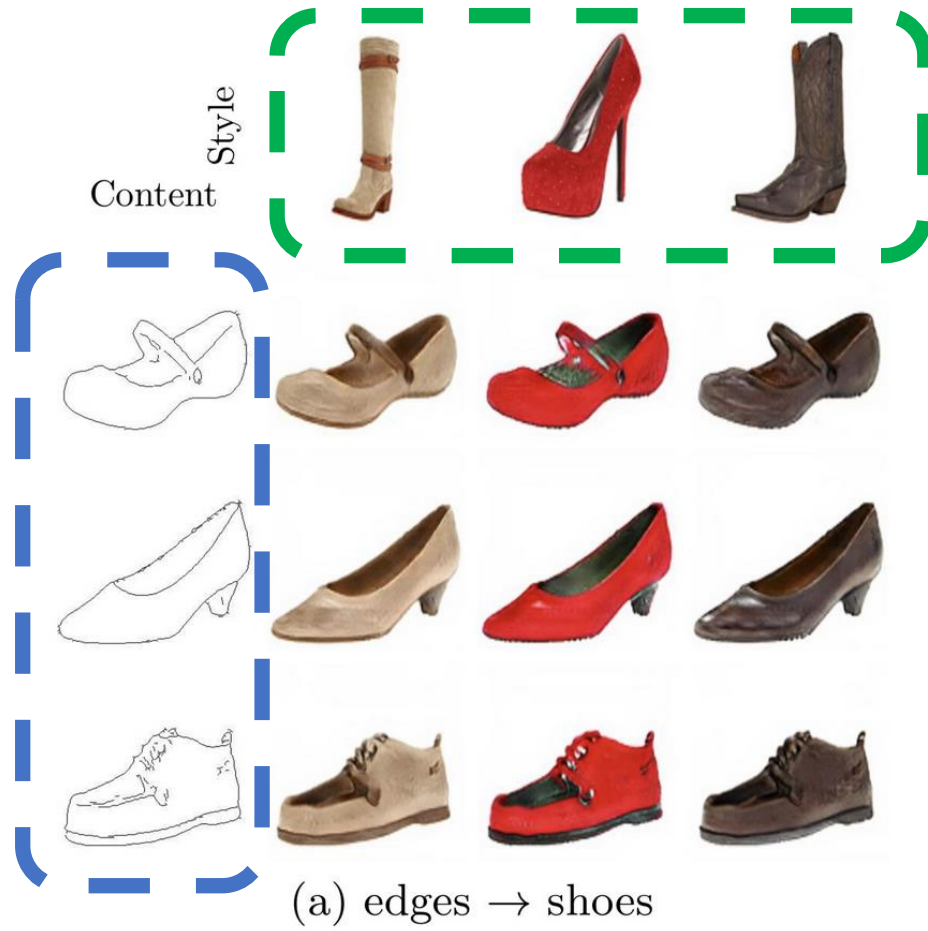
Generation / Synthesis



Karas et al, StyleGAN2, 2019

Synthesis with Style Transfer

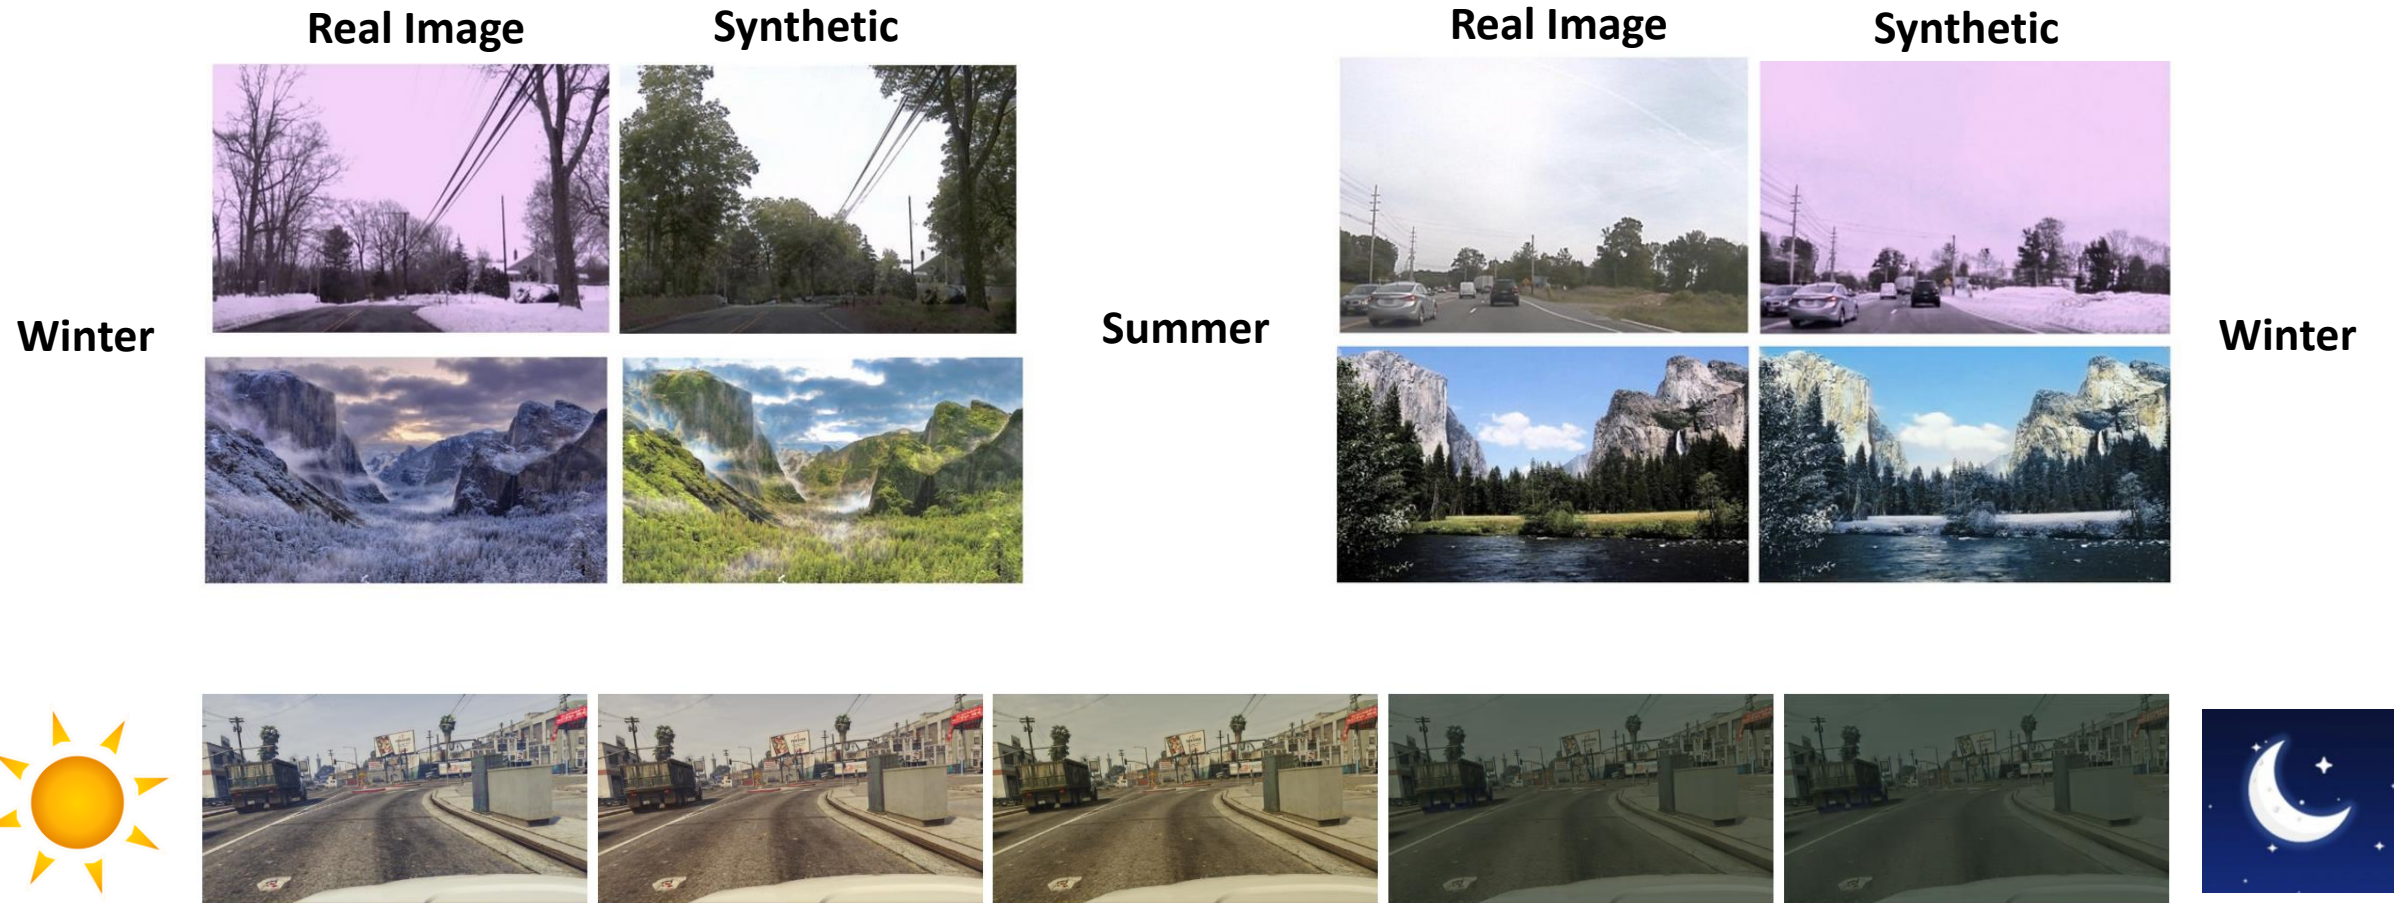
Example-guided Image Translation



Huang et al, Multimodal Unsupervised Image-to-Image Translation, 2018

Huang et al, Arbitrary Style Transfer in Real-time with Adaptive Instance Normalization, 2017

Style Transfer



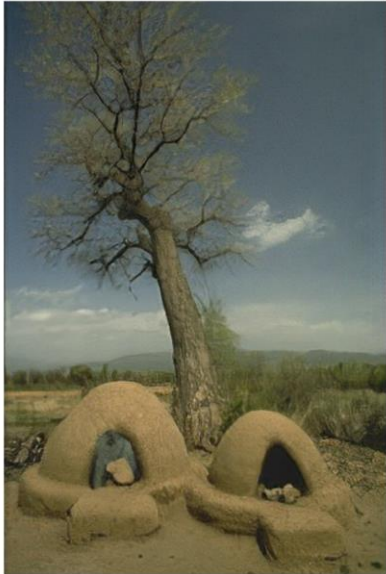
Huang et al, Multimodal Unsupervised Image-to-Image Translation, 2018
Huang et al, Arbitrary Style Transfer in Real-time with Adaptive Instance Normalization, 2017

Image Enhancement

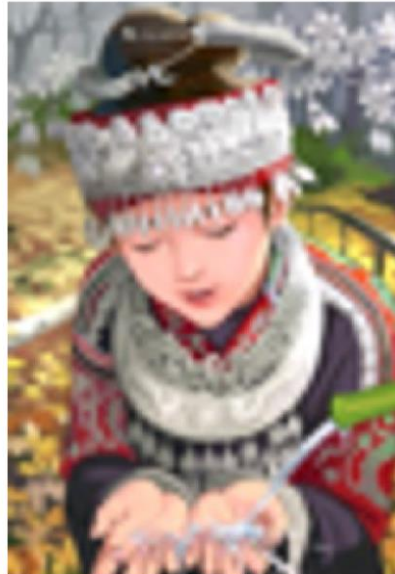
Down-sampled



GAN enhancement



Down-sampled



GAN enhancement



Down-sampled



GAN enhancement



Ledig et al, CVPR 2017

How do we learn from unlabeled data?

Next part:

Auto-Encoders (1)

Auto-Encoders (Part 1)

Neural Computation Course 2022

Konstantinos Kamnitsas

What we will learn about Auto-Encoders:

Part 1 (this part):

- What is the basic auto-encoder
- How to train auto-encoders
- What is an AE with a bottleneck layer and why to use it
- What features do they learn

Part 2:

- What can they be used for (dimensionality reduction, clustering, pretraining,...)
- What AEs are not good at and why

Supervised learning:

From observations (x) to labels (y), given labeled data

$$x \in \mathcal{X}$$

Observed variable: Because we see this data, both at training and at testing.
E.g. an image, or the vector representing a data point.

$$y \in \mathcal{Y}$$

Labeled variable: Used for variables for which a human created manual annotations. Usually the prediction **target** for **supervised** learning.

Observation:



Label:

“zero”



“one”

Unsupervised Learning:

From observations (x) to latent variables (z), without labeled data

$$x \in \mathcal{X}$$

Observed variable: Because we see this data, both at training and at testing.
E.g. an image, or the vector representing a data point.

$$z \in \mathcal{Z}$$

Latent variable: We do not know this information.
Given x, we **have to infer it**.



z: label of digit, thickness, slope ...
x: image of digit



z: content, lighting angle, zoom ...
x: the photo



"**Staff** makes you feel like **family** and the **quality** is out of this world."



"Amazing jeweler, great **prices**, geat **service!!!**"



"The **atmosphere** of the **showroom** is nice, you will feel relaxed, never pressured."

z: sentiment, vocabulary, grammar skills ...
x: written review

Example of latent (vector) representation

Observed variable

x = Input
vector

$x_1 =$



$$h = \begin{bmatrix} x_1^{(1)} \\ x_1^{(2)} \\ \vdots \\ x_1^{(w*h)} \end{bmatrix}$$

$x_2 =$



$$= \begin{bmatrix} x_2^{(1)} \\ x_2^{(2)} \\ \vdots \\ x_2^{(w*h)} \end{bmatrix}$$

Latent variable

$$z = \begin{bmatrix} z^{(1)} \\ z^{(2)} \\ z^{(3)} \\ z^{(4)} \\ \vdots \\ z^{(v)} \end{bmatrix} = \begin{bmatrix} \text{"blonde"} \\ \text{"smiling"} \\ \text{"looks_left"} \\ \text{"looks_right"} \\ \vdots \\ \text{"wears_suit"} \end{bmatrix}$$

$$z_1 = \begin{bmatrix} z_1^{(1)} \\ z_1^{(2)} \\ z_1^{(3)} \\ z_1^{(4)} \\ \vdots \\ z_1^{(v)} \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}$$

$$z_2 = \begin{bmatrix} z_2^{(1)} \\ z_2^{(2)} \\ z_2^{(3)} \\ z_2^{(4)} \\ \vdots \\ z_2^{(v)} \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 0 \\ 1 \\ \vdots \\ 1 \end{bmatrix}$$

How to learn useful latent representation Z of data?



$$f_{\phi} : \mathcal{X} \rightarrow \mathcal{Z}$$

...given only input samples x ...




$$z = \begin{bmatrix} \text{"blonde"} \\ \text{"smiling"} \\ \text{"looks_left"} \\ \text{"looks_right"} \\ \vdots \\ \text{"wears_suit"} \end{bmatrix}$$

$$z = \begin{bmatrix} \text{"slope"} \\ \text{"thickness"} \\ \vdots \\ \text{"is_circular"} \end{bmatrix}$$

How to learn useful latent representation (code) of data?


Supervised Classifier: $\mathcal{X} \rightarrow \mathcal{Y}$

Given for training:
Input & label (x, y)



Unsupervised Encoder: $\mathcal{X} \rightarrow \mathcal{Z}$

Given for training:
Only input samples x

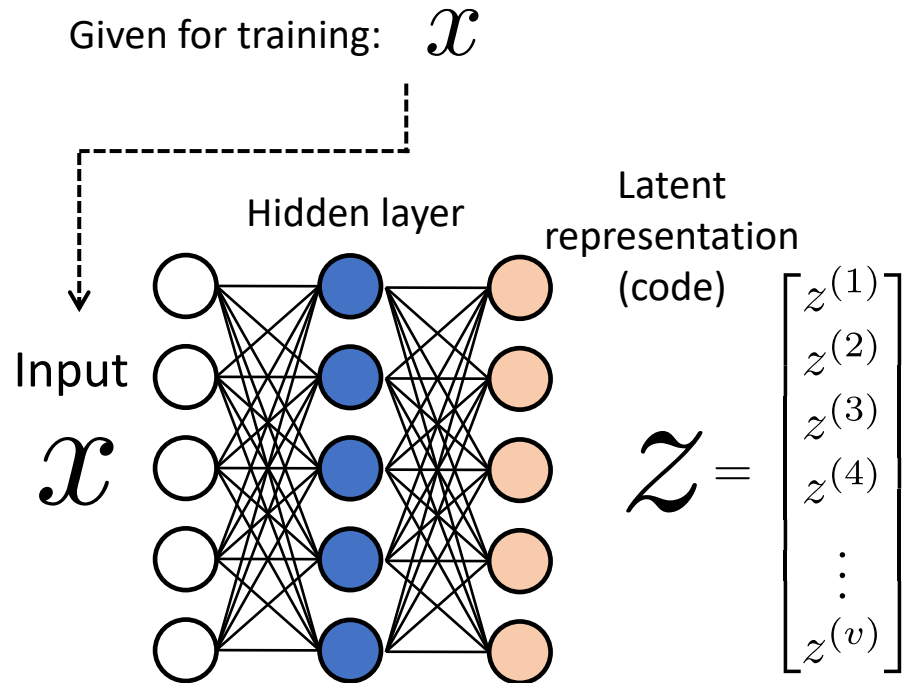


How do we train
such an encoder?

What Loss?

How to learn useful latent representation (code) of data?

Unsupervised Encoder: $f_\phi : \mathcal{X} \rightarrow \mathcal{Z}$



Main idea:

*What **property** do we require for latent representation \mathcal{Z} ?*



To represent \mathcal{X}



To **preserve useful info** about x



How can we enforce that?!?!?




Ensure we can also learn

$$g_\theta : \mathcal{Z} \rightarrow \mathcal{X}$$

How to learn useful latent representation (code) of data?

Unsupervised Auto-Encoder (AE):

Given for training: \mathcal{X}



Parameters: ϕ θ

Encoder:

Takes input x and encodes it to $z = f_\phi(x)$

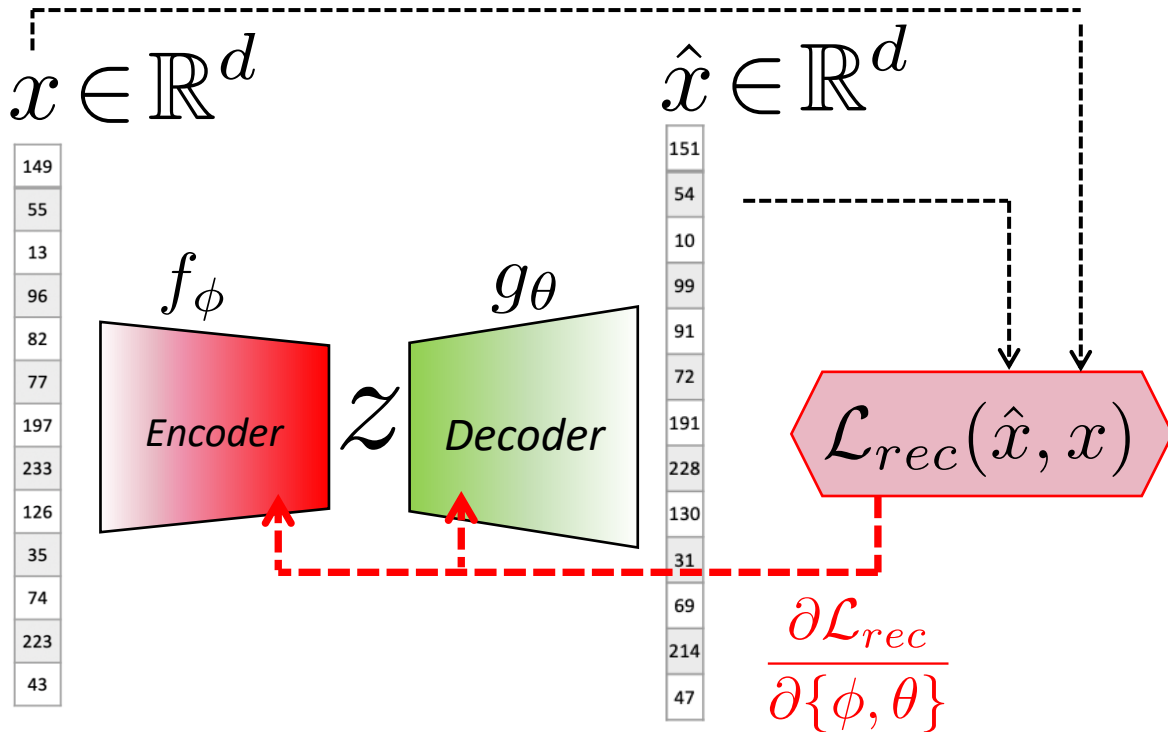
Decoder:

Takes code z and decodes it to re-construct $\hat{x} = g_\theta(f_\phi(x))$


Result:

Good reconstruction ($\hat{x} \approx x$) is possible only if code Z preserves info about X

Training loss



$$\phi', \theta' = \arg \min_{\phi, \theta} \mathcal{L}_{rec}$$



Reconstruction Loss
a.k.a. **Mean Squared Error:**

$$\begin{aligned} \mathcal{L}_{rec} &= \frac{1}{d} \sum_{j=1}^d (x^{(j)} - \hat{x}^{(j)})^2 \\ &= \frac{1}{d} \sum_{j=1}^d \left(x^{(j)} - g_\theta^{(j)}(f_\phi(x)) \right)^2 \end{aligned}$$

Red arrows point to $g_\theta^{(j)}$ and $f_\phi(x)$ in the second equation, indicating the parameters being updated during training.

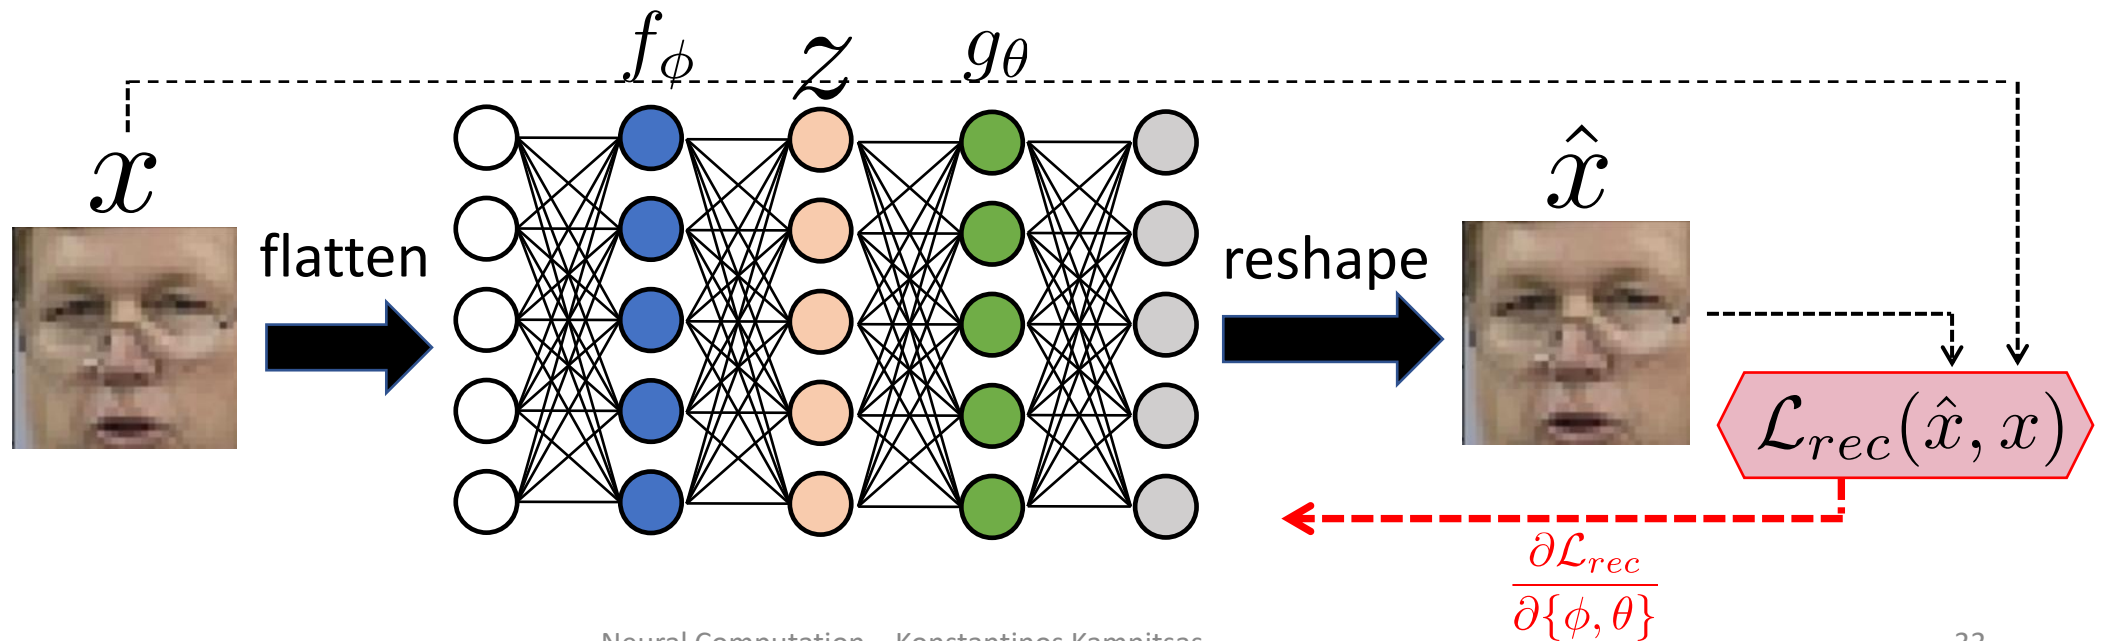
Only one Global Optimum:

$$\begin{aligned} \mathcal{L}_{rec} = 0 &\Leftrightarrow \hat{x}^{(j)} - x^{(j)} = 0 \quad \forall j \\ &\Leftrightarrow \boxed{\hat{x} = x} \end{aligned}$$

Learning to reconstruct

The loss is minimized when: $\hat{x} = x \Rightarrow \underbrace{g_{\theta}(f_{\phi}(x))}_{\text{identity function}} = x$

The Auto-Encoder tries to learn the **identity function**!

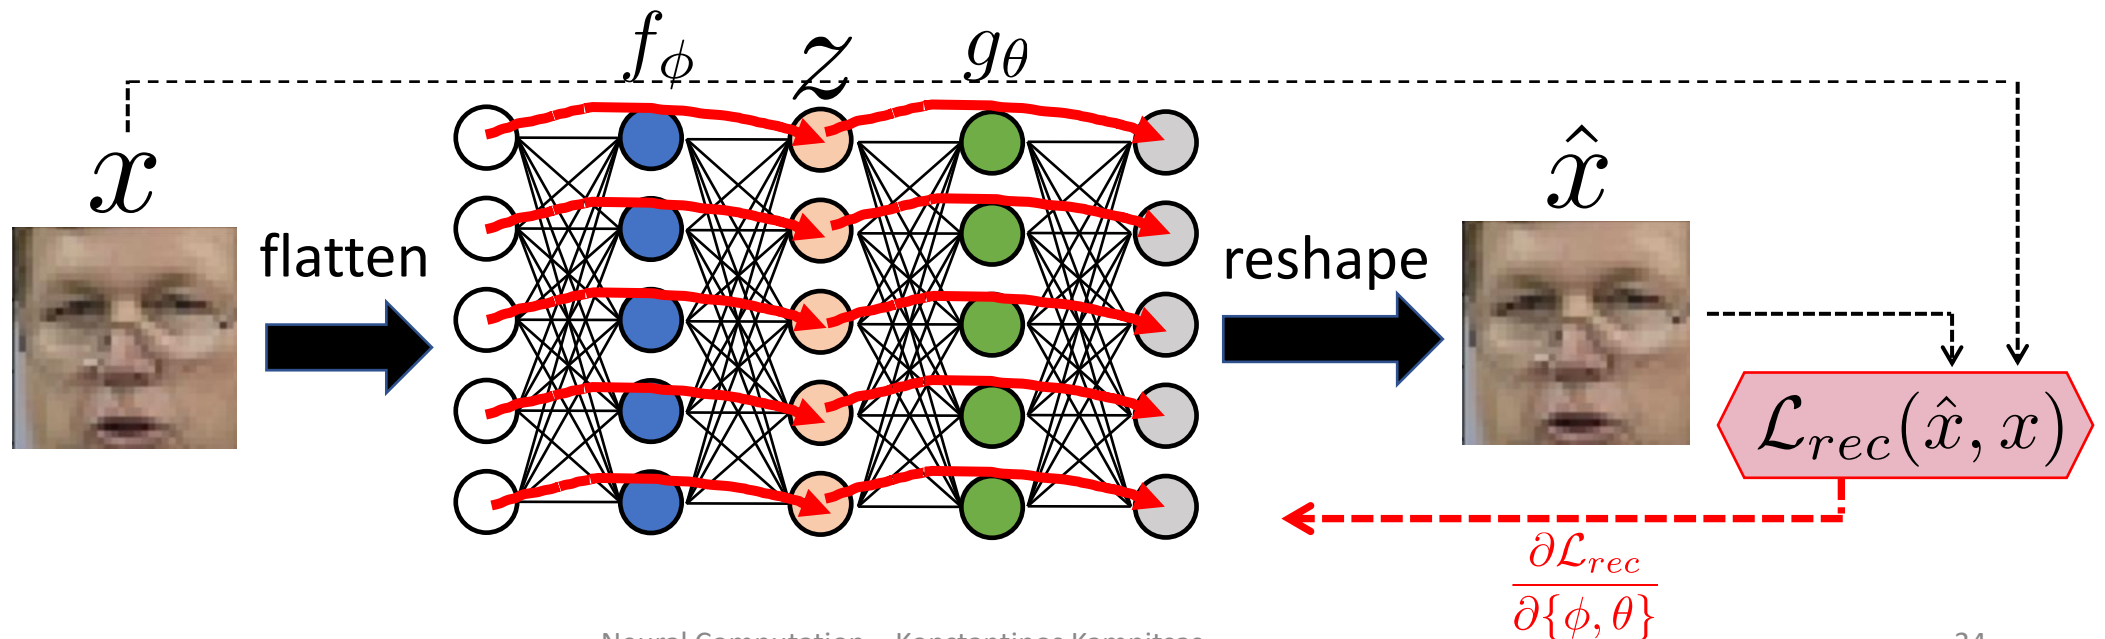


Problem: Trivial Solution / Learning a useless function

The loss is minimized when: $\hat{x} = x \Rightarrow \underbrace{g_{\theta}(f_{\phi}(x))}_{\text{identity function}} = x$

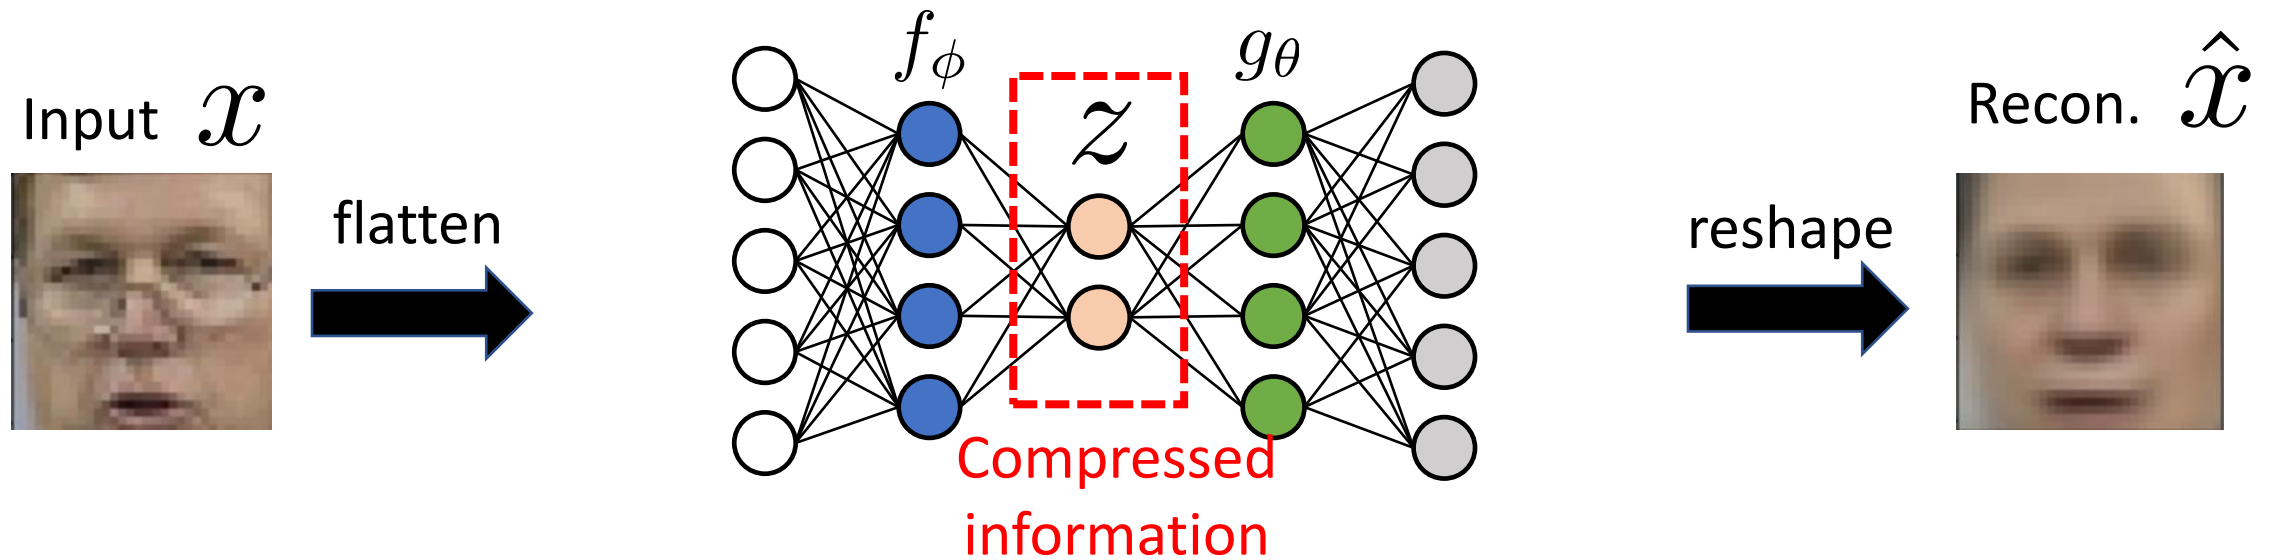
The Auto-Encoder tries to learn the **identity function**!

Exists ***trivial solution (and useless model)***: $z = f_{\phi}(x) = x$
 $\hat{x} = g_{\theta}(z) = z$



Solution: *Bottleneck* layer

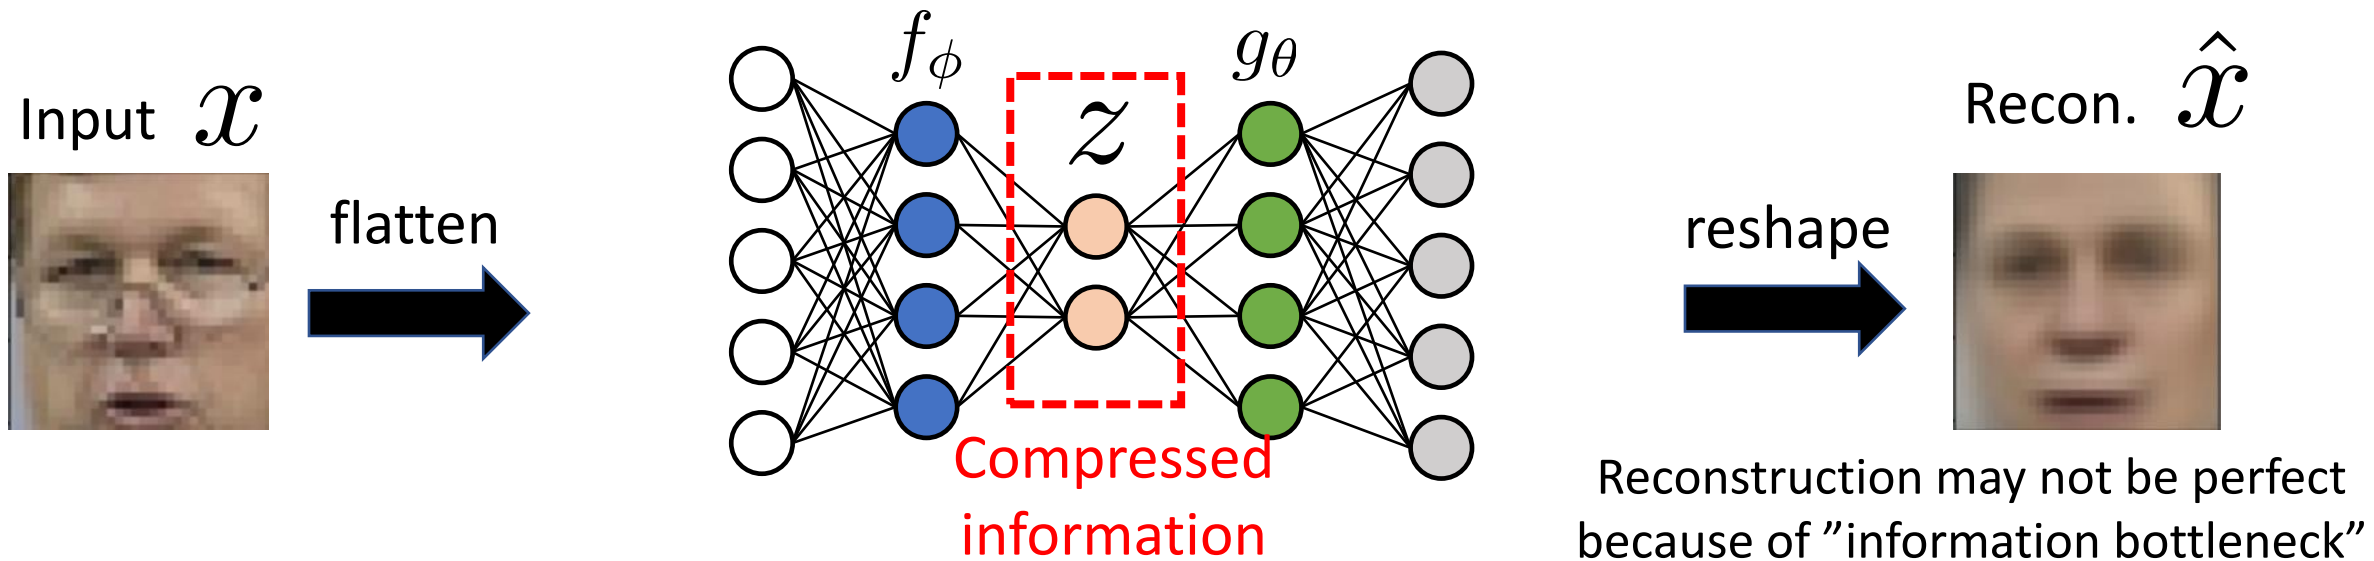
$$f_{\phi} : \mathcal{X} \in \mathbb{R}^d \rightarrow \mathcal{Z} \in \mathbb{R}^v, \text{ where } v < d$$



"Standard/basic" Auto-Encoder

$$f_{\phi} : \mathcal{X} \in \mathbb{R}^d \rightarrow \mathcal{Z} \in \mathbb{R}^v, \text{ where } v < d$$

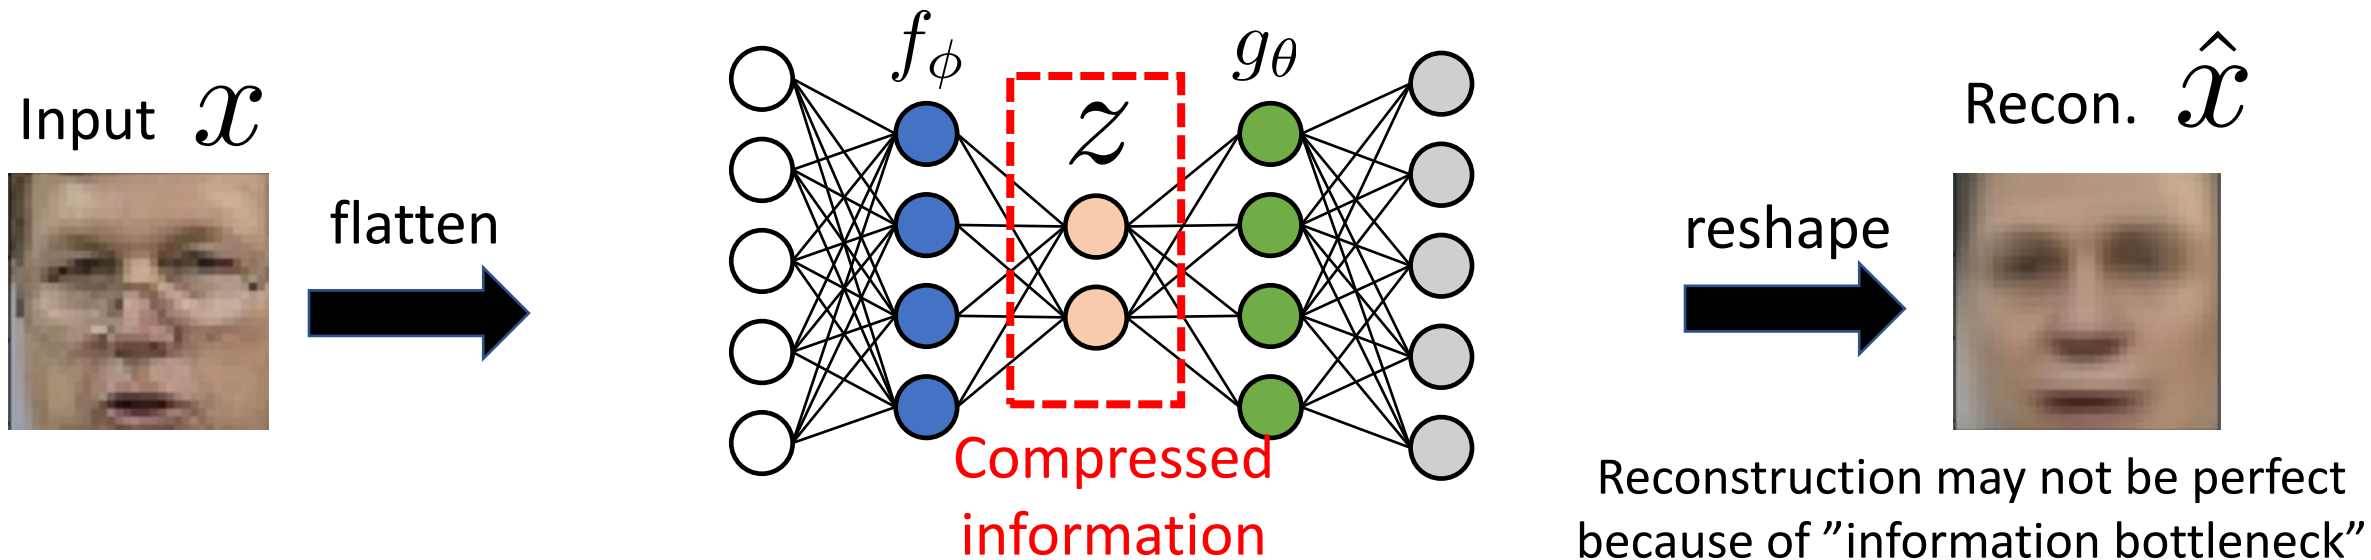
*With "standard/basic" Auto-Encoder,
in this class we refer to
AE with Bottleneck*



"Standard/basic" Auto-Encoder

$$f_{\phi} : \mathcal{X} \in \mathbb{R}^d \rightarrow \mathcal{Z} \in \mathbb{R}^v, \text{ where } v < d$$

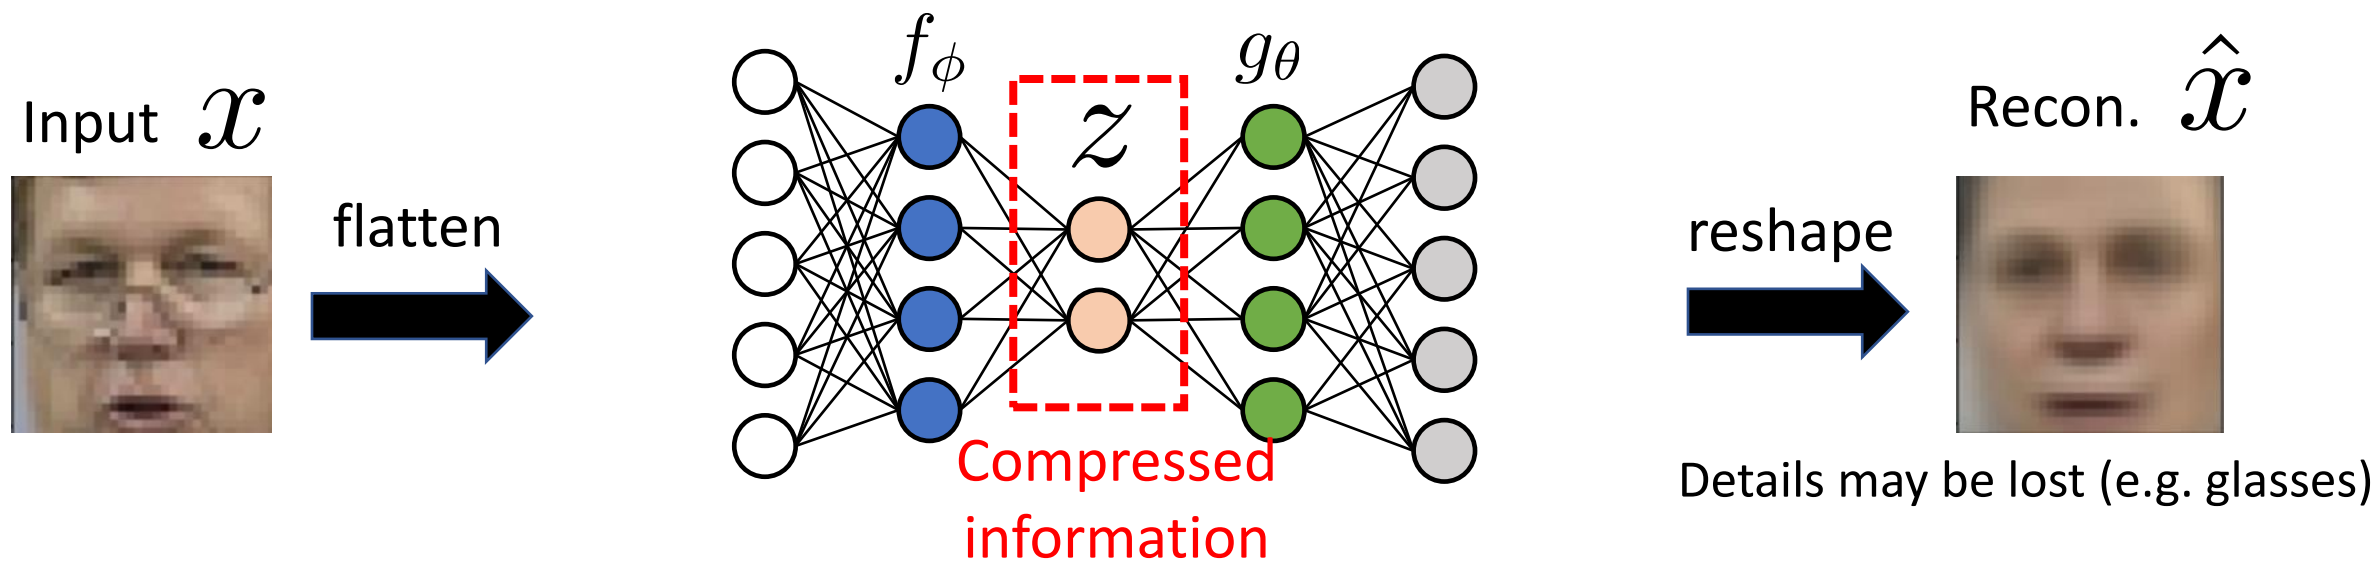
*With "standard/basic" Auto-Encoder,
in this class we refer to
AE with Bottleneck*



***What do you think would
the AE learn to encode in \mathcal{Z} ?***

Compression forces encoding the most prominent features

$$f_{\phi} : \mathcal{X} \in \mathbb{R}^d \rightarrow \mathcal{Z} \in \mathbb{R}^v, \text{ where } v < d$$



Reconstruction loss penalizes wrong pixel intensities. Therefore:


Encoder usually learns to encode features that **explain intensities of as many pixels as possible**.

Usually these are **“high level”** features, as often called in deep learning:

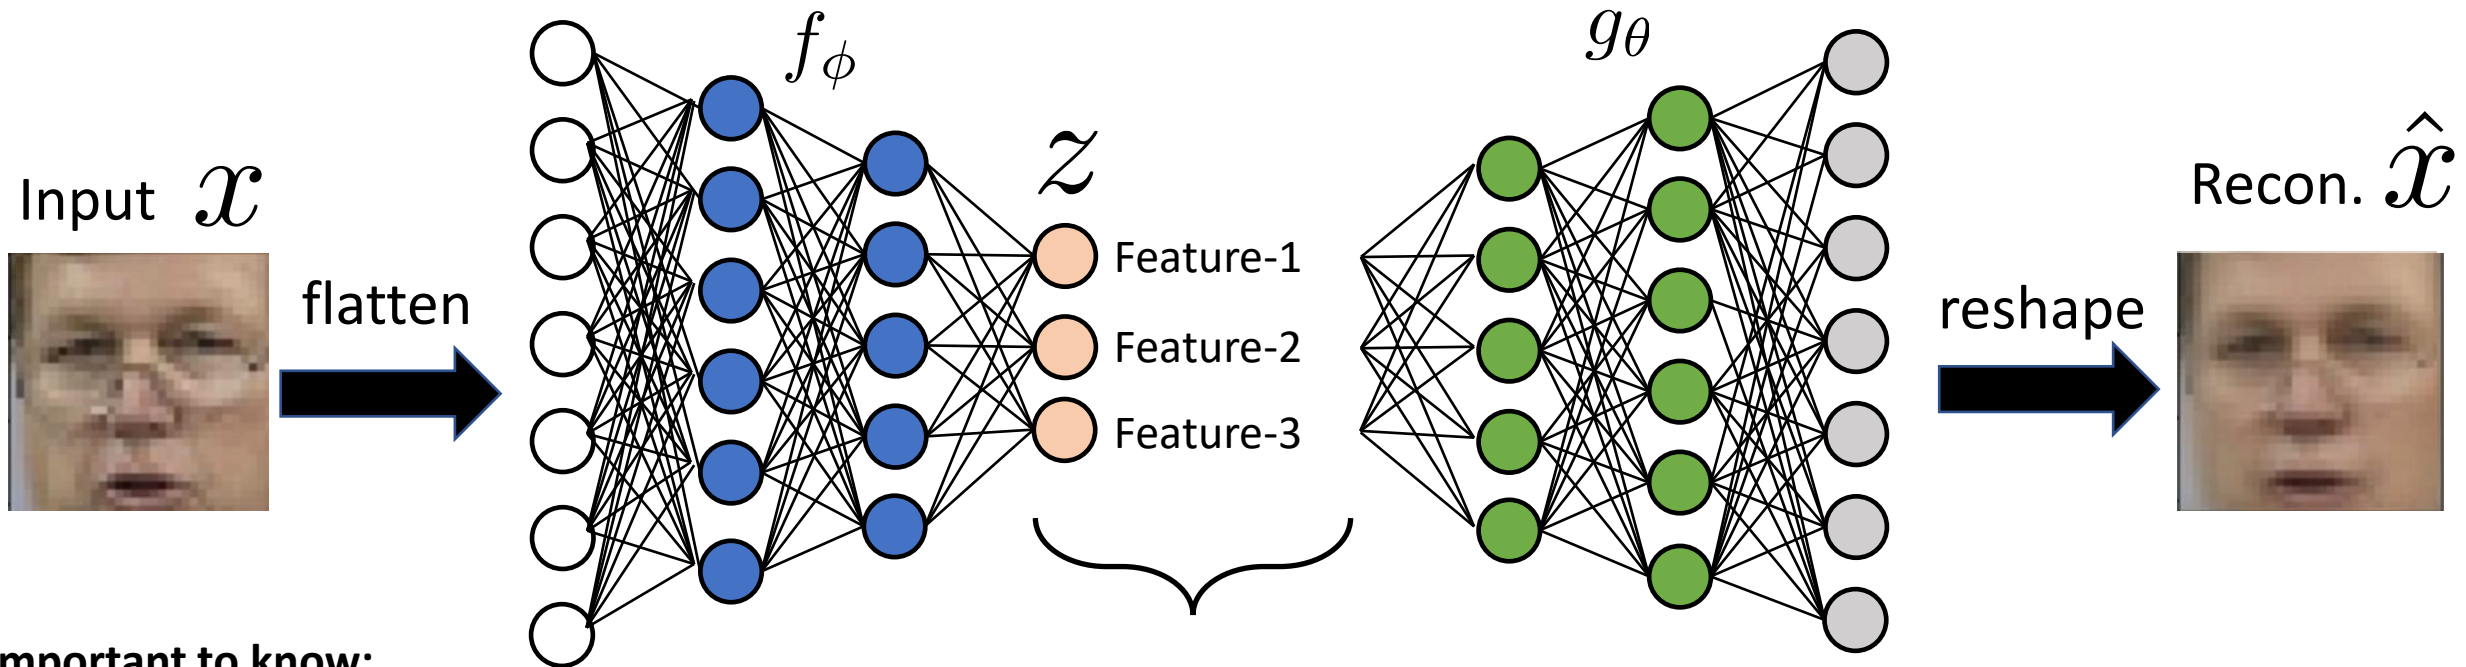
E.g. here: Skin color (most pixels), location and size of eyes, mouth, nose (dark areas), type of hair, clothing...

Wider bottleneck: Better reconstruction, less compression

$$f_{\phi} : \mathcal{X} \in \mathbb{R}^d \rightarrow \mathcal{Z} \in \mathbb{R}^v, \text{ where } v < d$$





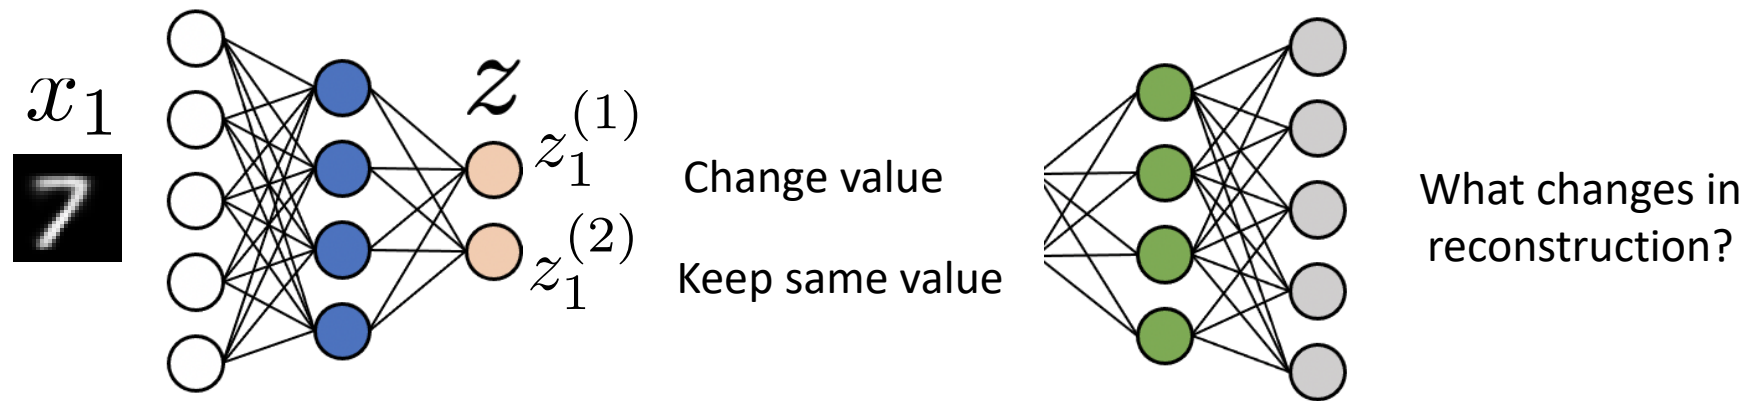
Bottleneck forces the AE to learn to encode the **few most important features** of data. But what are they?




Important to know:

- **We do not control** what features are learned
- Re-training with different random seeds **may learn different features** due to randomness of SGD!
- **We do not know** what features (attributes) of the input are learned
- We can find out by visual inspection: After an AE was trained, encode x and decode it after changing only 1 feature. Then, check what changed in the reconstruction!

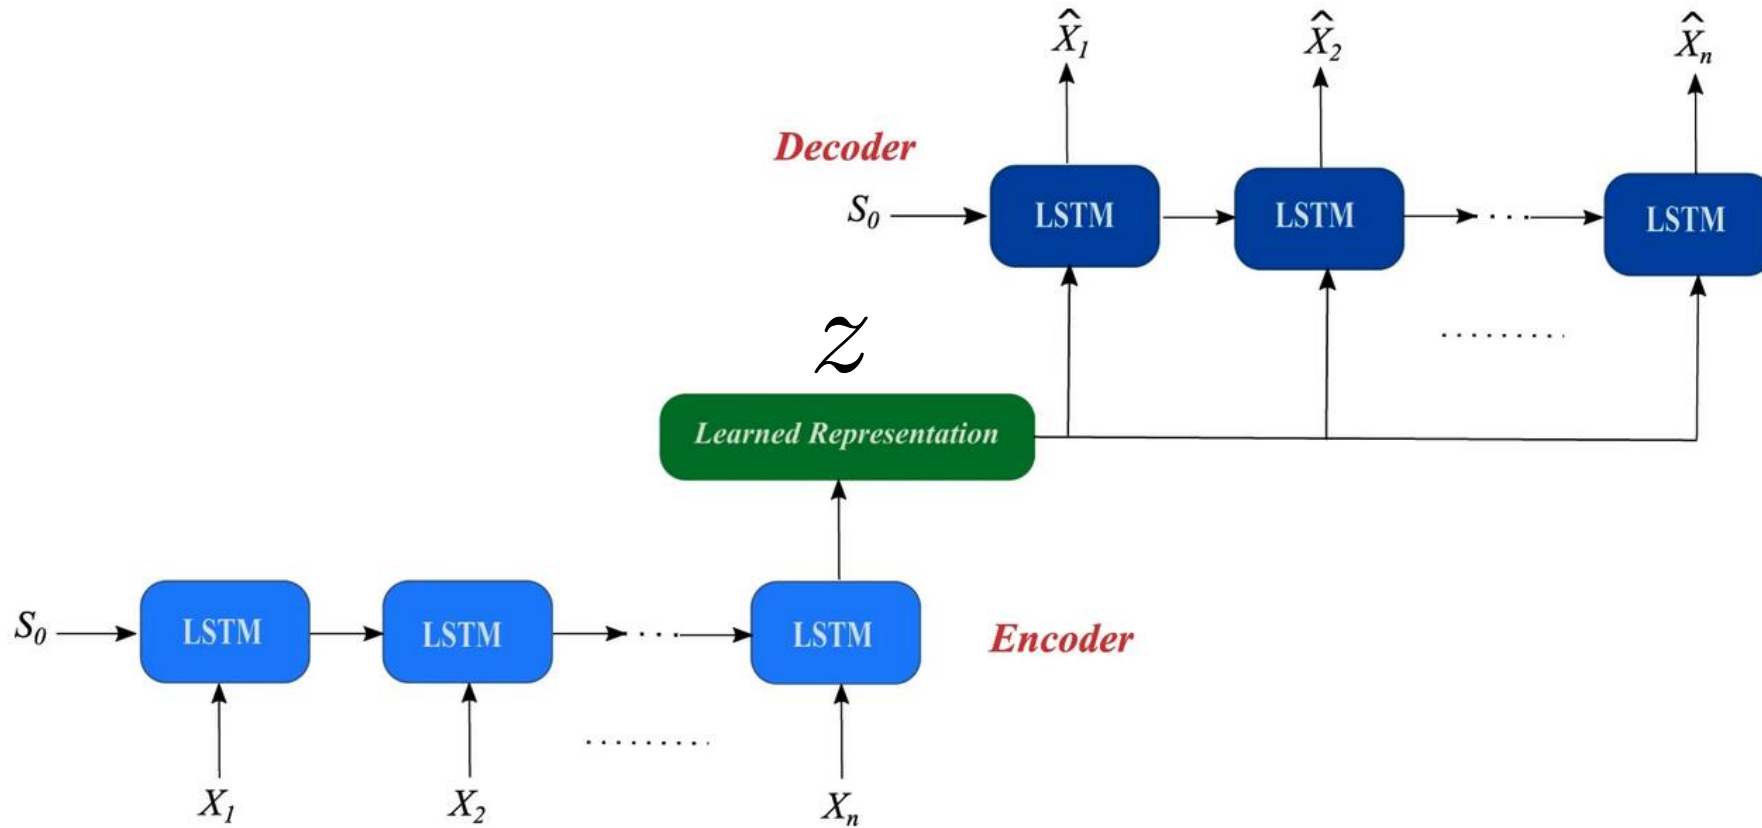
Investigating what the features are by visual inspection



Convolutional Auto-Encoders



RNN-based Auto-Encoders



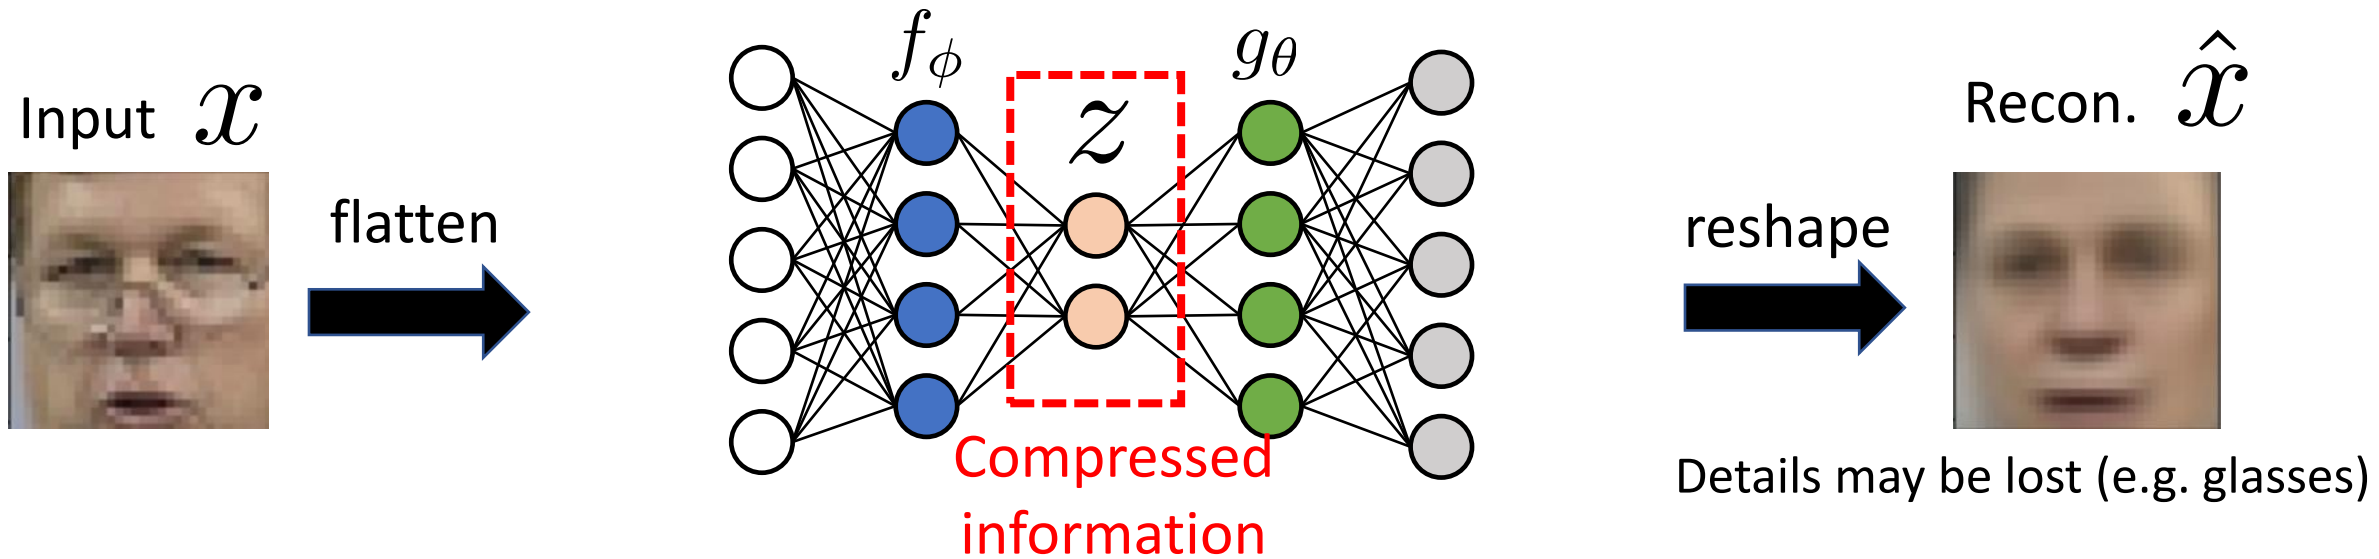
In this part:

- What is an Auto-Encoder
- What is the information bottleneck and why use it
- How to train an Auto-Encoder

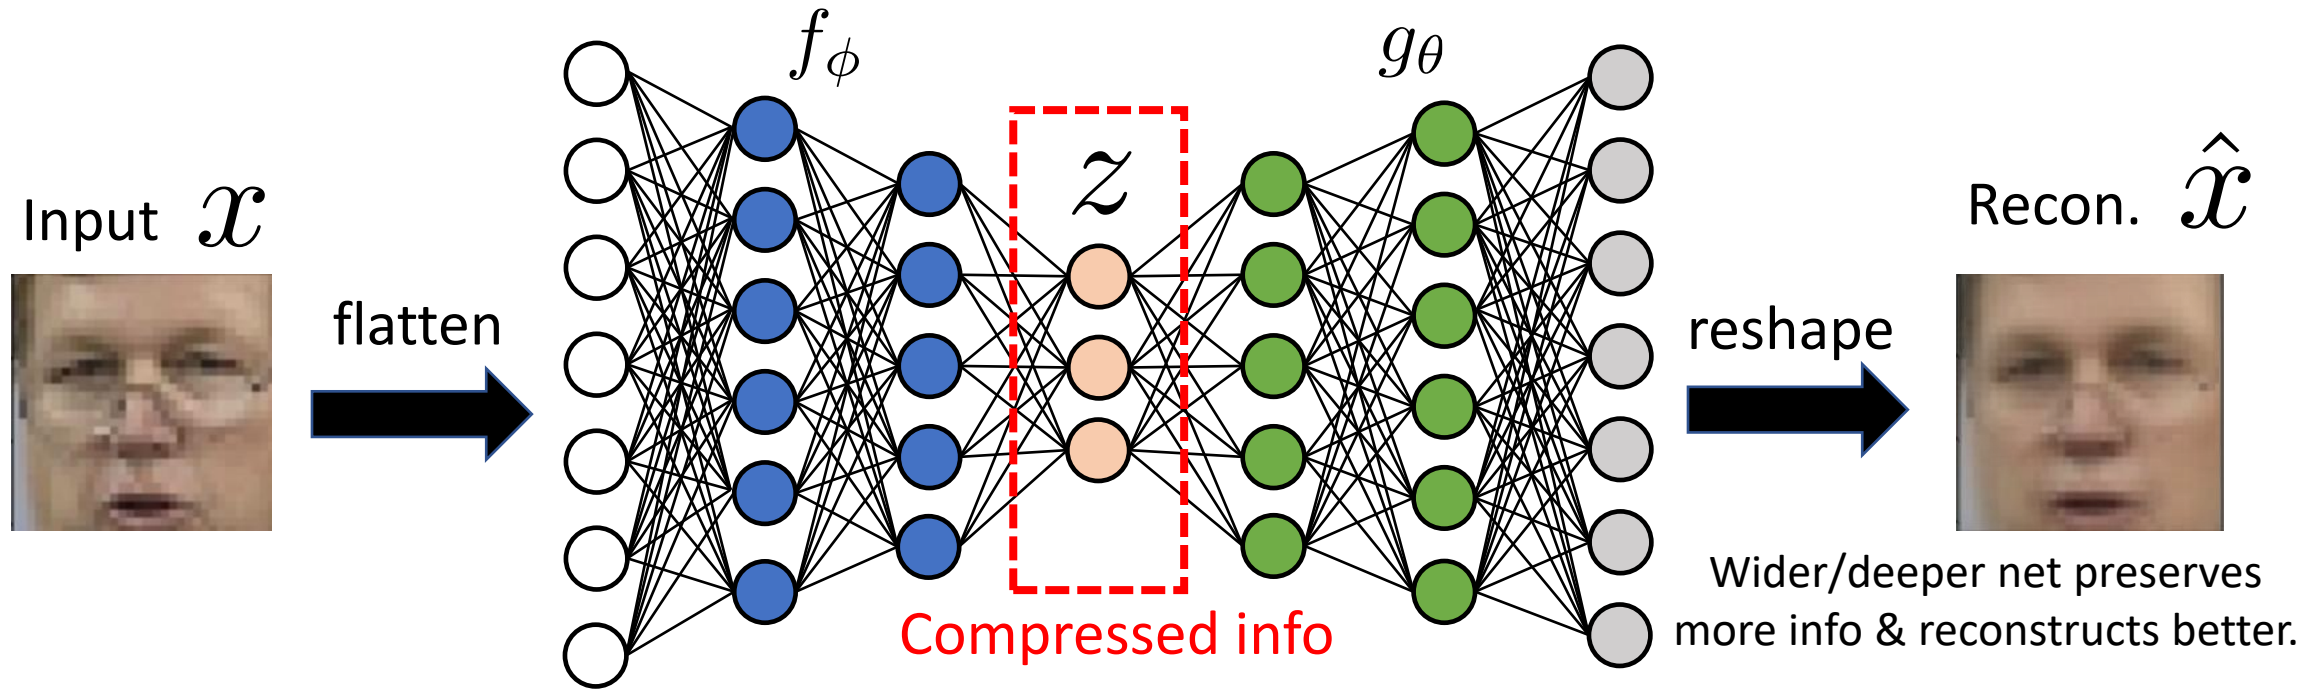
Next part:

- What can we **use AEs for?**
- What are AEs **not good for?**

AE with Bottleneck layer learns to perform Dimensionality Reduction / Compression via Unsupervised Learning



Wider bottleneck: Better reconstruction, less compression



But if too little compression \Rightarrow AE may not learn high level features
(towards encoding all pixels / identity func).

It's a **trade-off**.

Use of AEs for compression - Example application

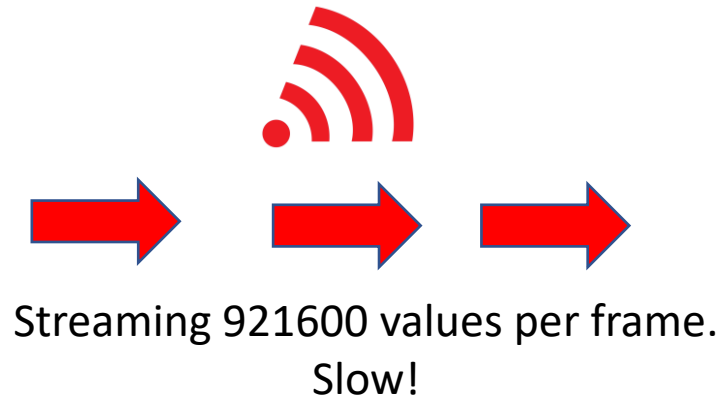
Streaming Company



720

1280

$1280 \times 720 = 921600$ pixels

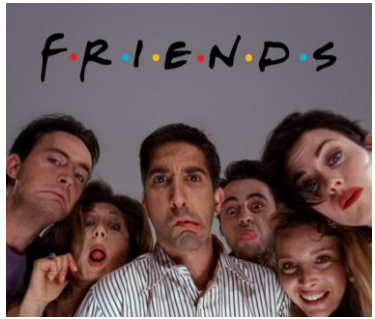


Customer



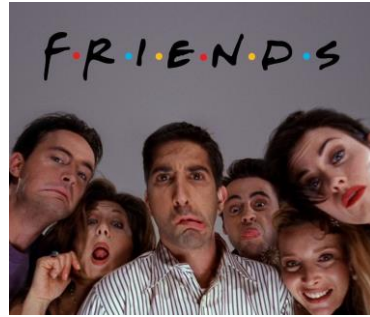

Use of AEs for compression - Example application

(1) Streaming company trains AE on the movie frames:



720


1280



$1280 \times 720 = 921600$ pixels

Use of AEs for compression - Example application

(1) Streaming company trains AE on the movie frames:




(2) Stores movie
encoded in Z
on own servers



Use of AEs for compression - Example application

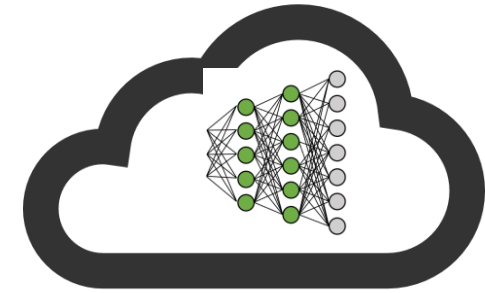
(1) Streaming company trains AE on the movie frames:



(2) Stores movie encoded in Z on own servers


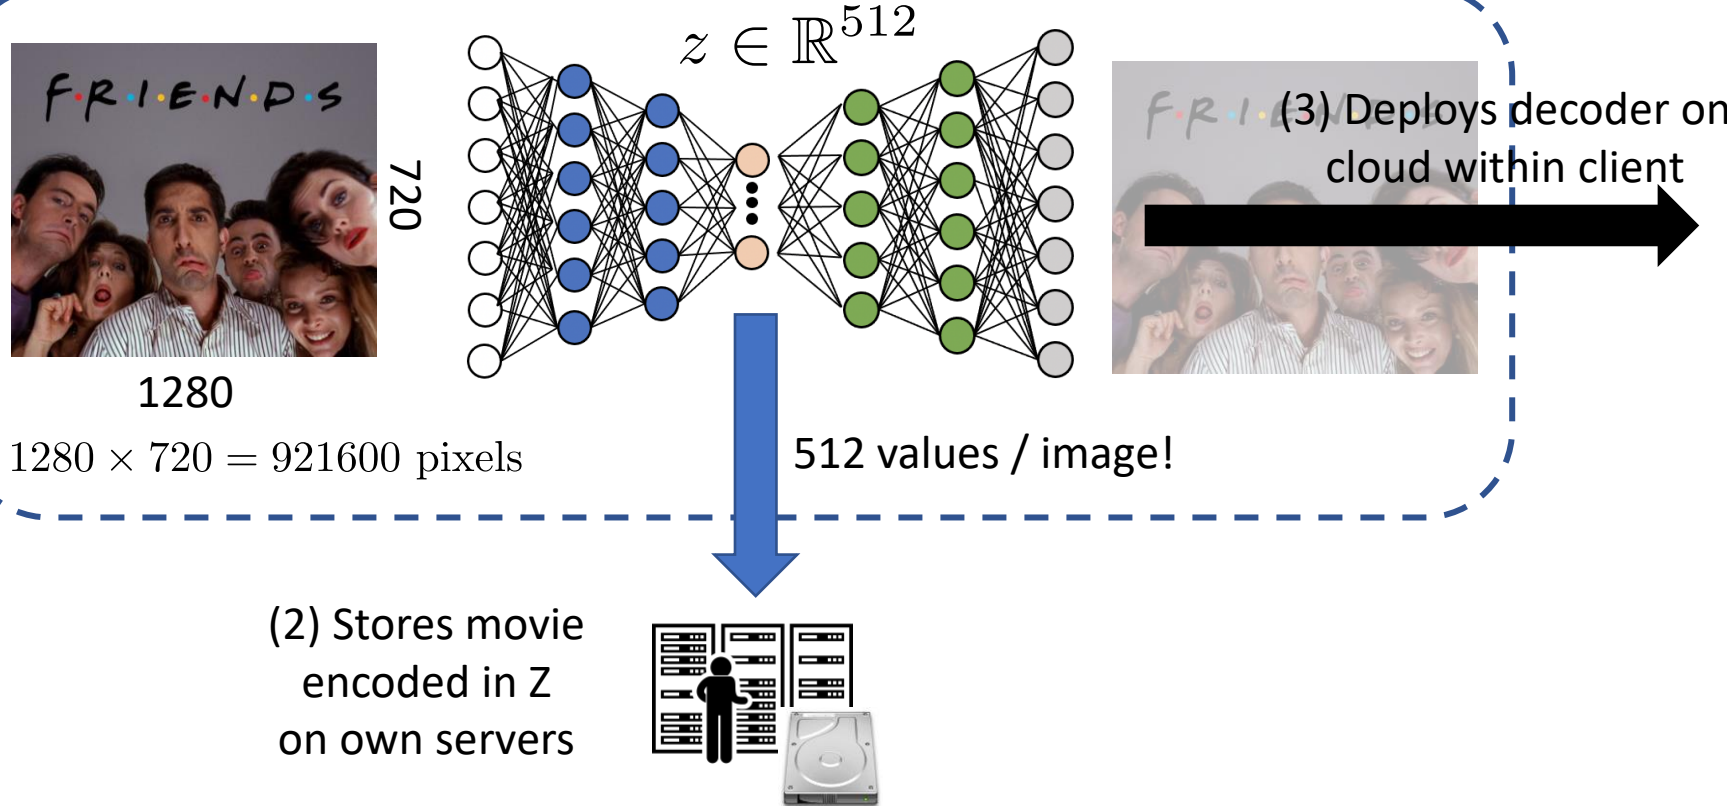


(3) Deploys decoder on cloud within client




Use of AEs for compression - Example application

(1) Streaming company trains AE on the movie frames:



Use of AEs for compression - Example application

(1) Streaming company trains AE on the movie frames:



(2) Stores movie encoded in Z on own servers




(5) Streaming movie:
Transfer only code z
(e.g. 512 values per frame)! Fast!




$z \in \mathbb{R}^{512}$

(3) Deploys decoder on cloud within client

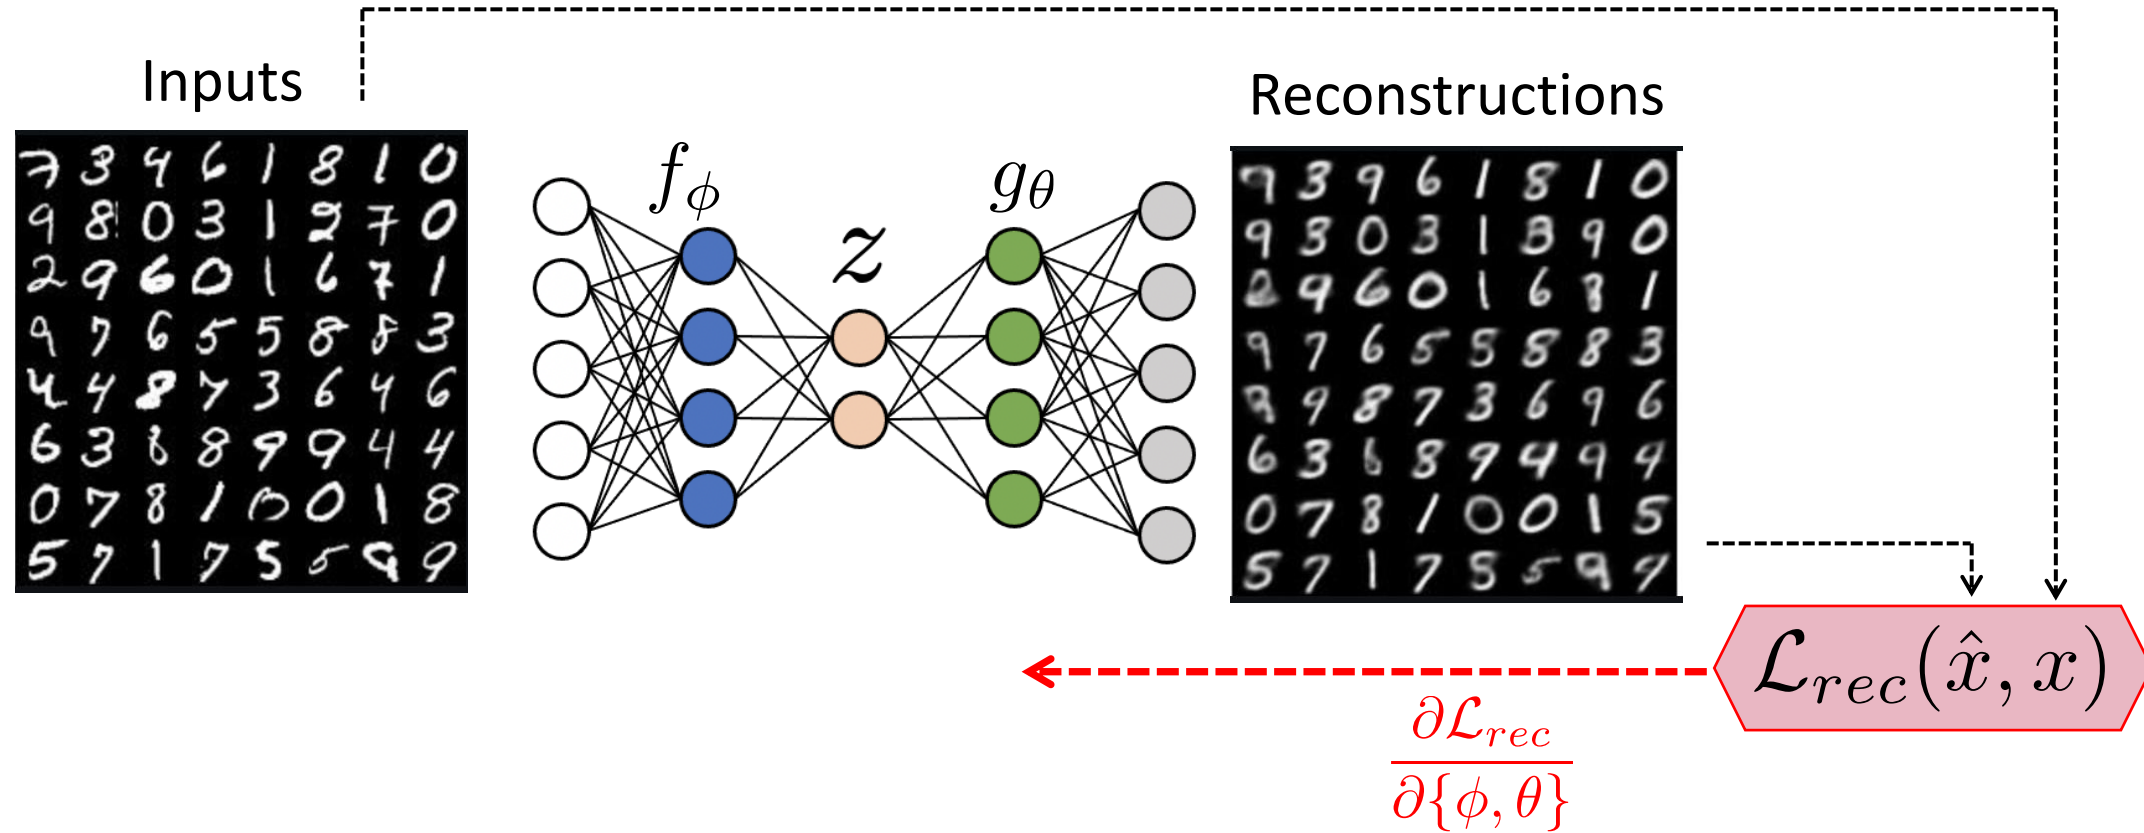


(4) Customer downloads decoder




(6) Decode $z \Rightarrow x$ on customer's laptop

AutoEncoders: Do they learn to encode “useful” info in Z?



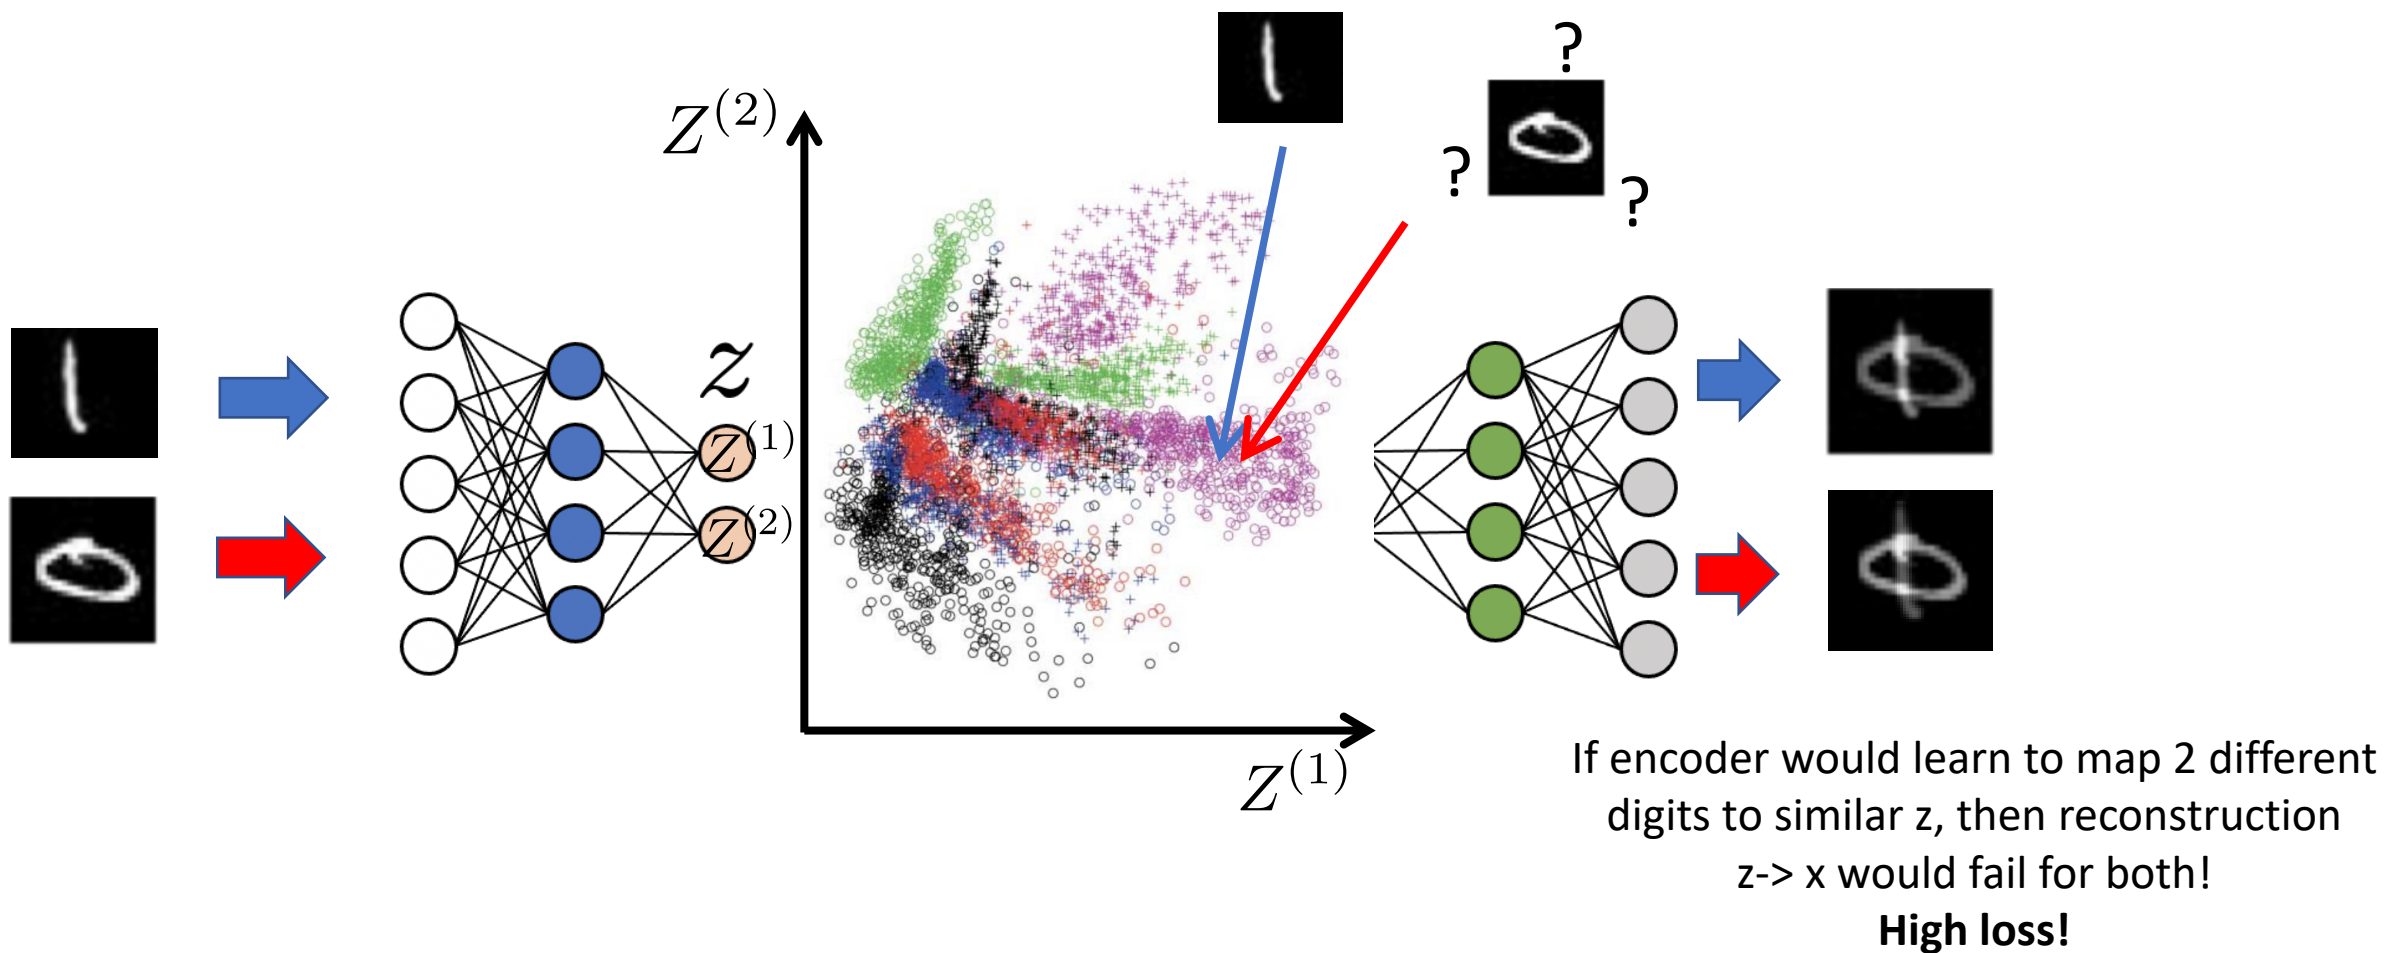
AEs can learn to cluster the data in unsupervised manner!

After training:

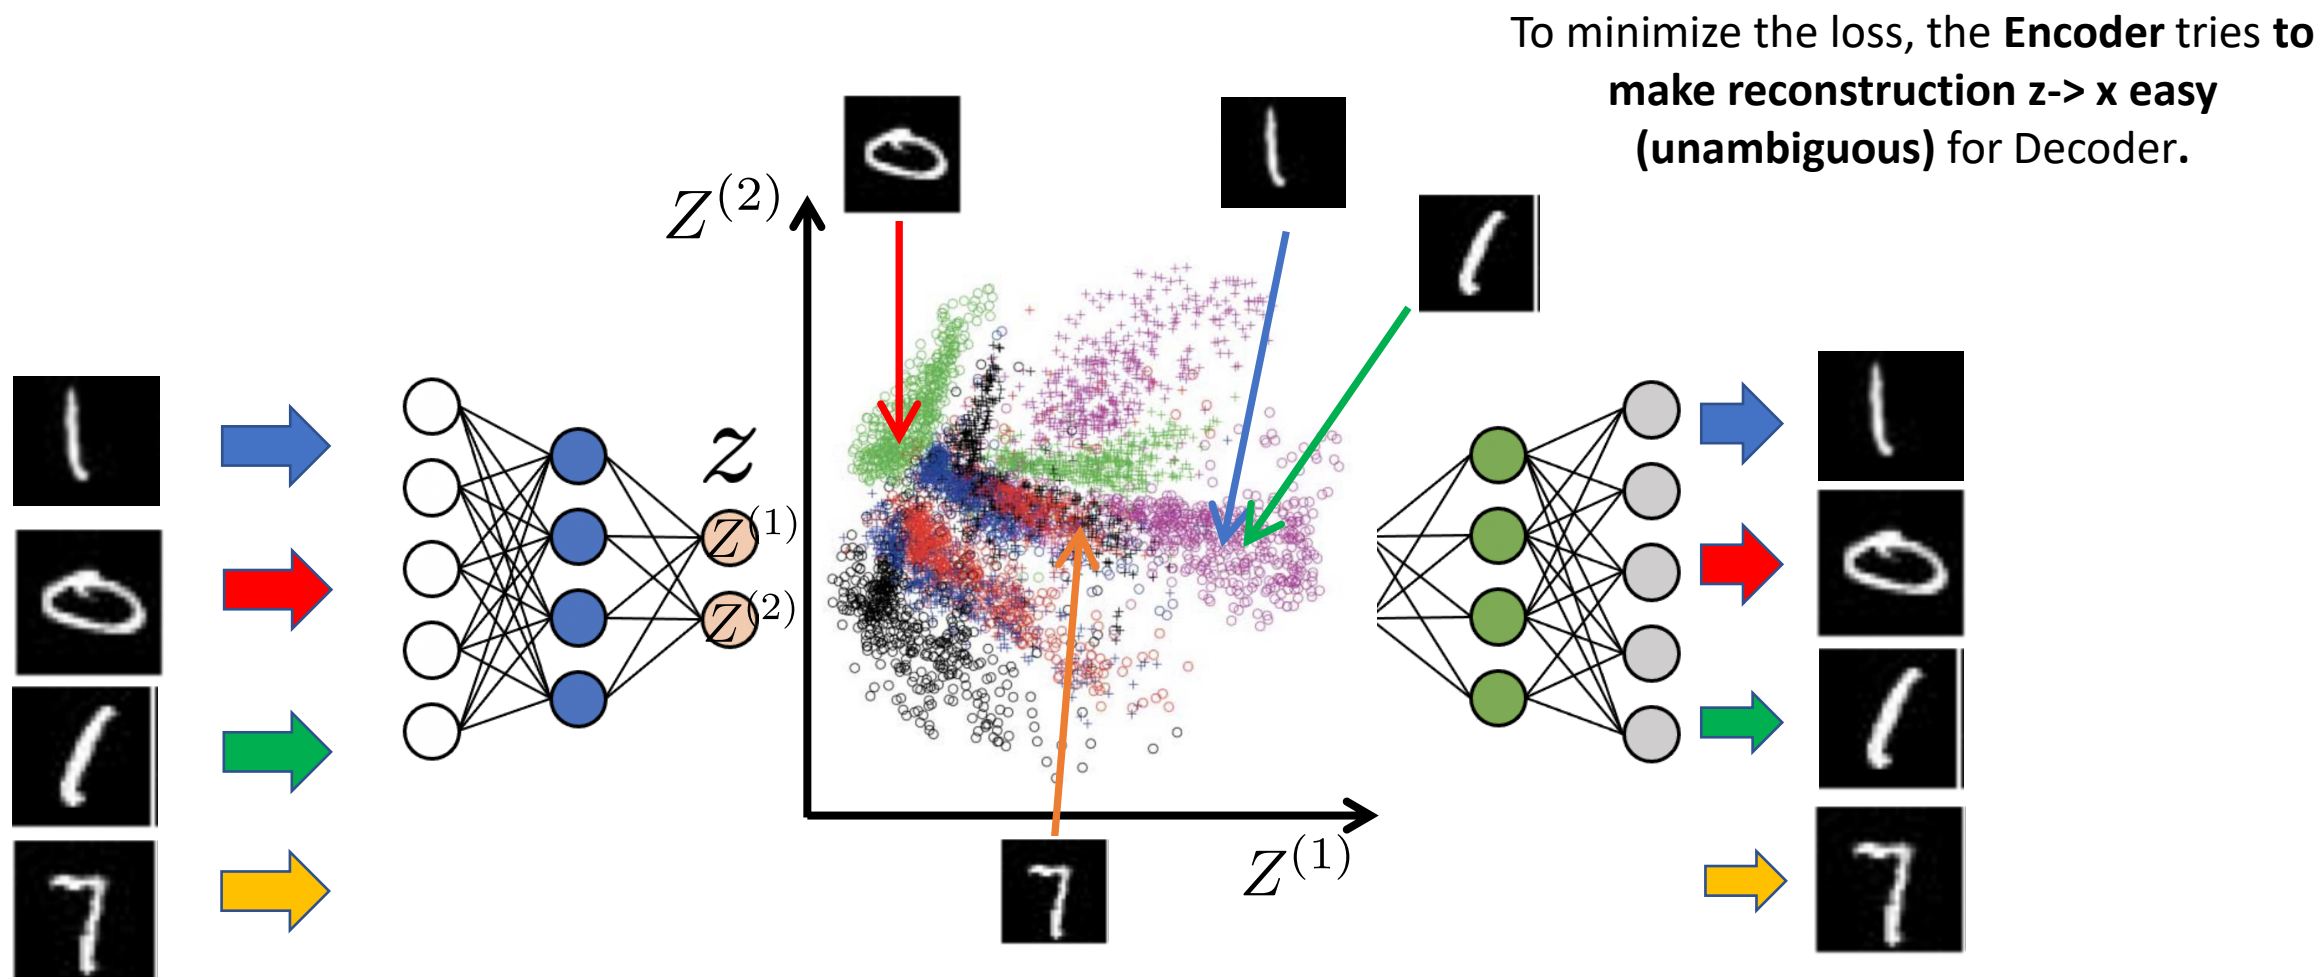


From: Hinton and Salakhutdinov, Reducing the Dimensionality of Data with Neural Networks, Science, 2006

Why learn to cluster data?



Why learn to cluster data?




Learn from unlabeled data, when labels are limited

Assume **our ultimate goal is to learn a Classifier with Supervised Learning.**


But number (N) of labeled data is **small**.

$$\mathbf{D} = (\mathbf{X}, \mathbf{Y}) = \{(x_i, y_i)\}_{i=1}^N$$




Potential overfit!
Whole image-to-label network
(many parameters) trained with
little data!


Pre-train with unlabeled data...



Trained Encoder of AE maps data X to space Z where they are **clustered**.
Throw away the decoder...




Attach **untrained** classifier on top of encoder
(commonly 1-2 layers)



Commonly shallow, 1-2 layers.
Therefore, it has very few parameters

(Approach 1)

Using the **limited labelled data**, train with **Supervised Learning ONLY** the classifier.
Keep parameters of encoder frozen.



Advantage:


Trains only the the small classifier with the limited labels. Therefore can avoid over-fit.

Disadvantage:

Encoder is not optimized for labelled data. May be suboptimal.

(Approach 2)

Using the limited labelled data, train with **Supervised Learning BOTH** the **encoder** and the **classifier** (usually for only few SGD iterations).



Advantage:

Encoder is optimized via labels, which "may" lead to better representation Z and results.

Disadvantage:

Possibility to overfit as all parameters are trained. Limited GD steps to avoid this.

Number of GD steps must be carefully decided on validation data to avoid overfit.


Can we generate (synthesize) new (not real) data with basic AE?

Can we generate (synthesize) new (not real) data with basic AE?

Assume an **already trained** Auto-Encoder

$$f : \mathcal{X} \rightarrow \mathcal{Z}$$

$$g : \mathcal{Z} \rightarrow \mathcal{X}$$




Is it a model appropriate for generating new data?
It is not trained for it..!


Can we generate (synthesize) new (not real) data with basic AE?

Step 1: Sample random z

E.g. With uniform probability between
[min, max] values z seen during training




Step 2: Decode



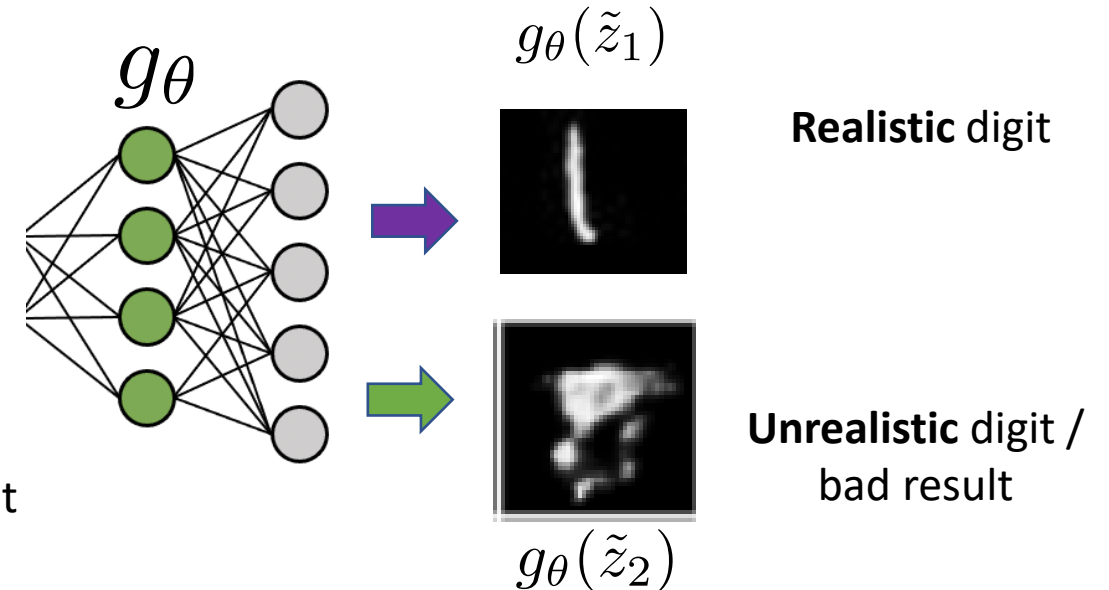
Problems generating new data with basic AE

Step 1: Sample random z

E.g. With uniform probability between [min, max] values z seen during training



Step 2: Decode



Problems:

- a) No “**real**” digits were encoded in that area during training. Hence these z values do not encode “realistic” digits.
- b) Decoder has not learned to decode such z values

Basic AEs are not appropriate for image generation.
Reconstruction loss does not train AE for generation.

We will see how **Generative Models (VAEs and GANs)** are trained appropriately for generation.

In this video:

- What can we use basic AEs for?
(e.g. compression, clustering, pretrain/initialize supervised model when labelled data are limited)
- What are “standard” AEs not good for?
(generation of new samples)

Next week:

- Generative Models
- Variational Auto-Encoders

For further reading (e.g. some more advanced AEs):

- Goodfellow, Bengion & Courville, Deep Learning, Chapter 14
<https://www.deeplearningbook.org/>

(However only the material our own slides, videos & tutorials will be assessed)

Thank you very much for your attention